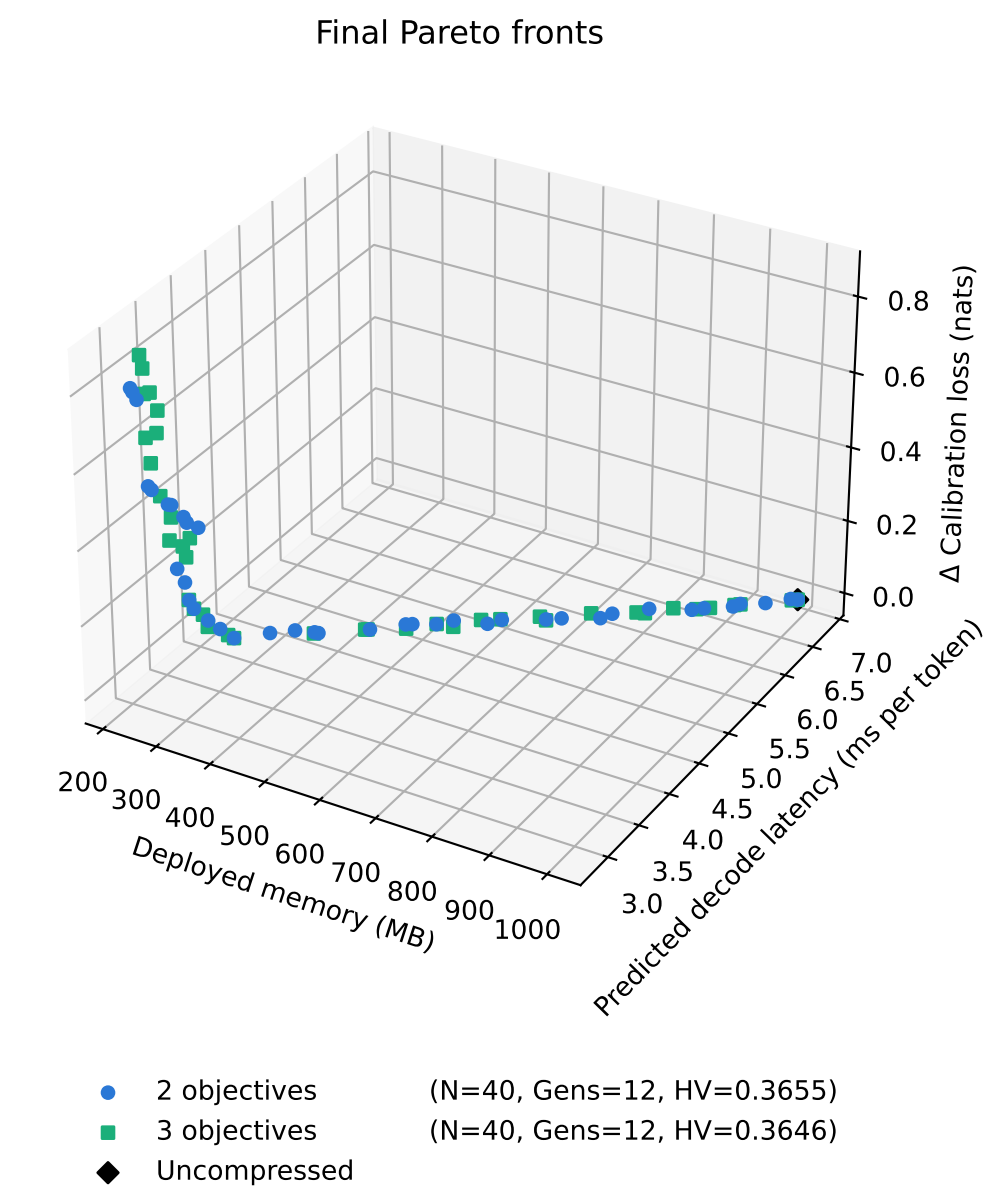
Final Pareto fronts
Deployed memory (MB)
200
300
400
500
600
700
800
900
1000
Predicted decode latency (ms per token)
3.0
3.5
4.0
4.5
5.0
5.5
6.0
6.5
7.0
Δ Calibration loss (nats)
0.0
0.2
0.4
0.6
0.8
2 objectives (N=40, Gens=12, HV=0.3655)
3 objectives (N=40, Gens=12, HV=0.3646)
Uncompressed