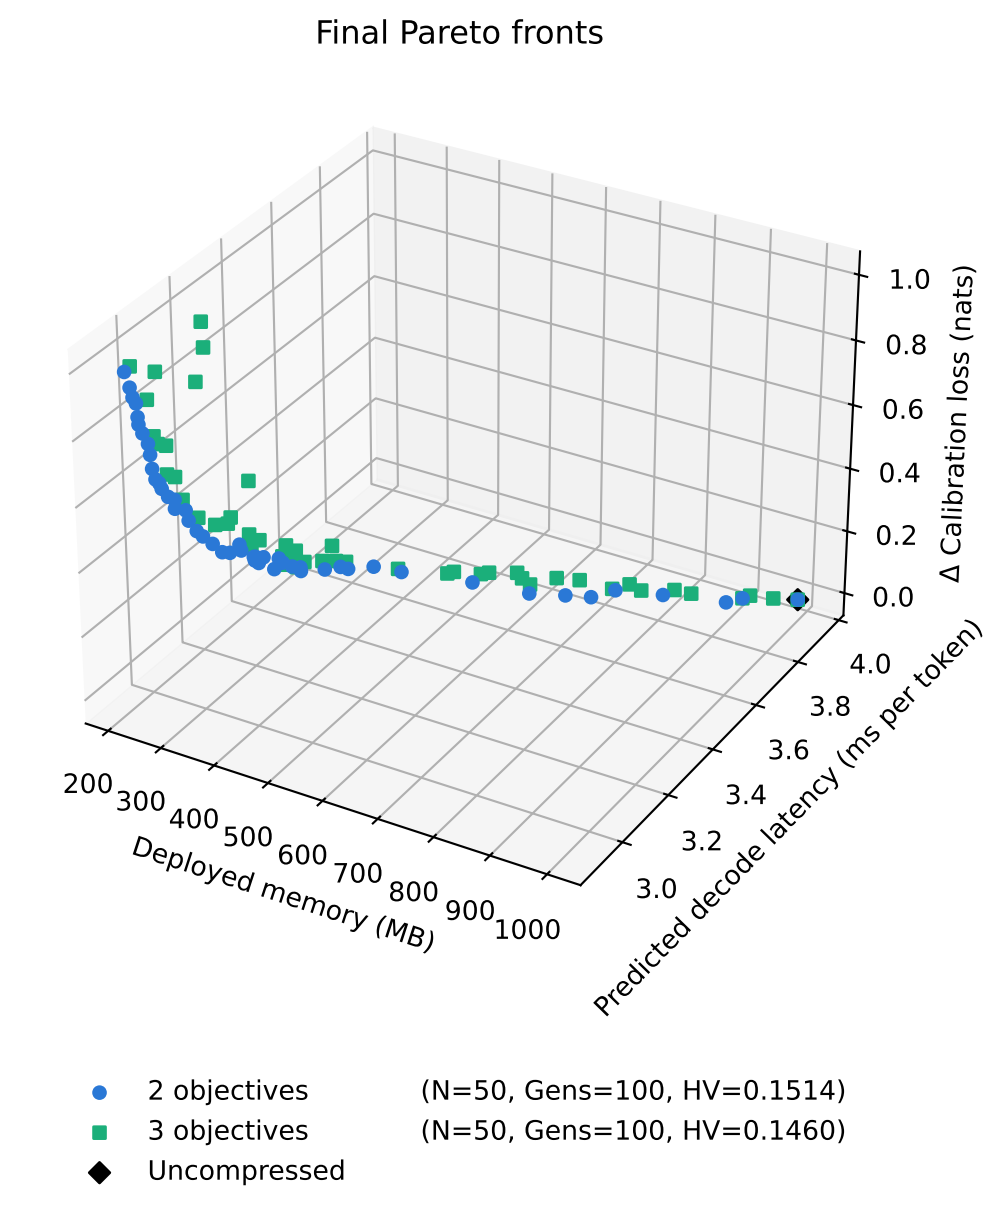
Final Pareto fronts
Deployed memory (MB)
200
300
400
500
600
700
800
900
1000
Predicted decode latency (ms per token)
3.0
3.2
3.4
3.6
3.8
4.0
Δ Calibration loss (nats)
0.0
0.2
0.4
0.6
0.8
1.0
2 objectives (N=50, Gens=100, HV=0.1514)
3 objectives (N=50, Gens=100, HV=0.1460)
Uncompressed