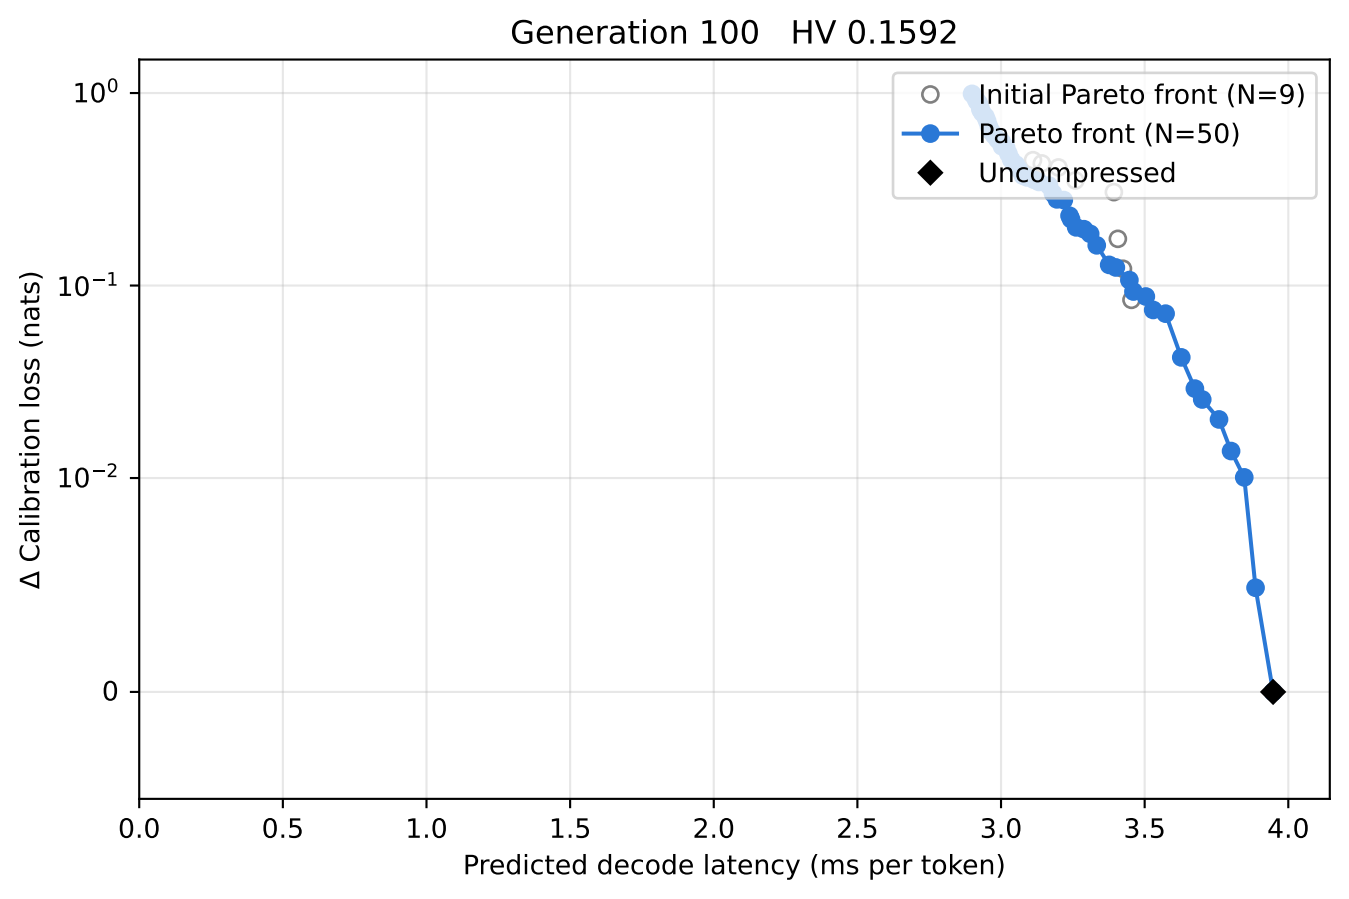
Generation 100 HV 0.1592
Initial Pareto front (N=9)
Pareto front (N=50)
Uncompressed
Δ Calibration loss (nats)
10^{0}
10^{-1}
10^{-2}
0
0.0
0.5
1.0
1.5
2.0
2.5
3.0
3.5
4.0
Predicted decode latency (ms per token)