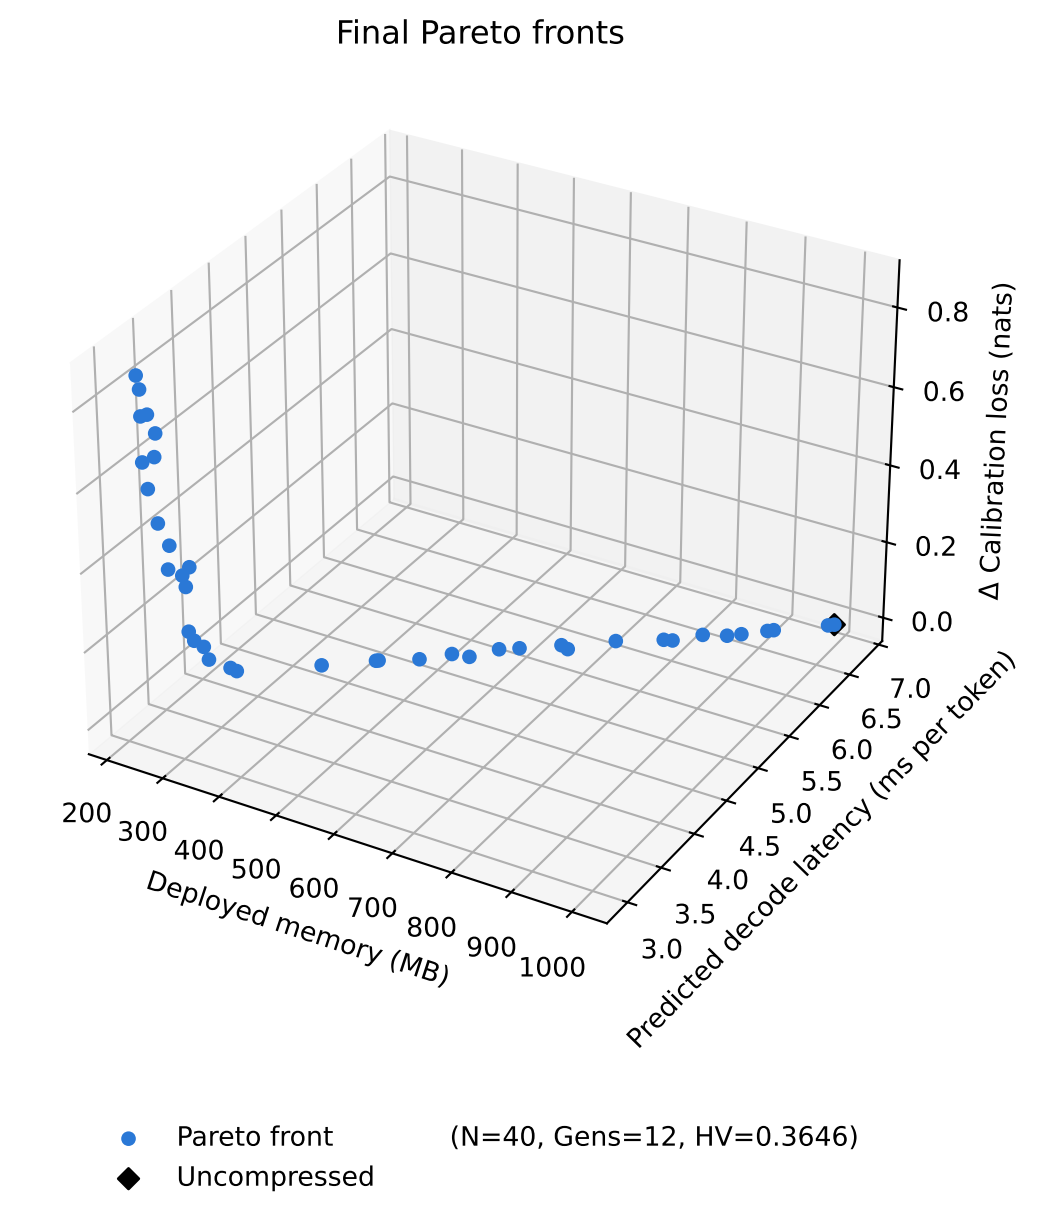
Final Pareto fronts
Δ Calibration loss (nats)
0.0
0.2
0.4
0.6
0.8
Predicted decode latency (ms per token)
3.0
3.5
4.0
4.5
5.0
5.5
6.0
6.5
7.0
Deployed memory (MB)
200
300
400
500
600
700
800
900
1000
Pareto front
(N=40, Gens=12, HV=0.3646)
Uncompressed