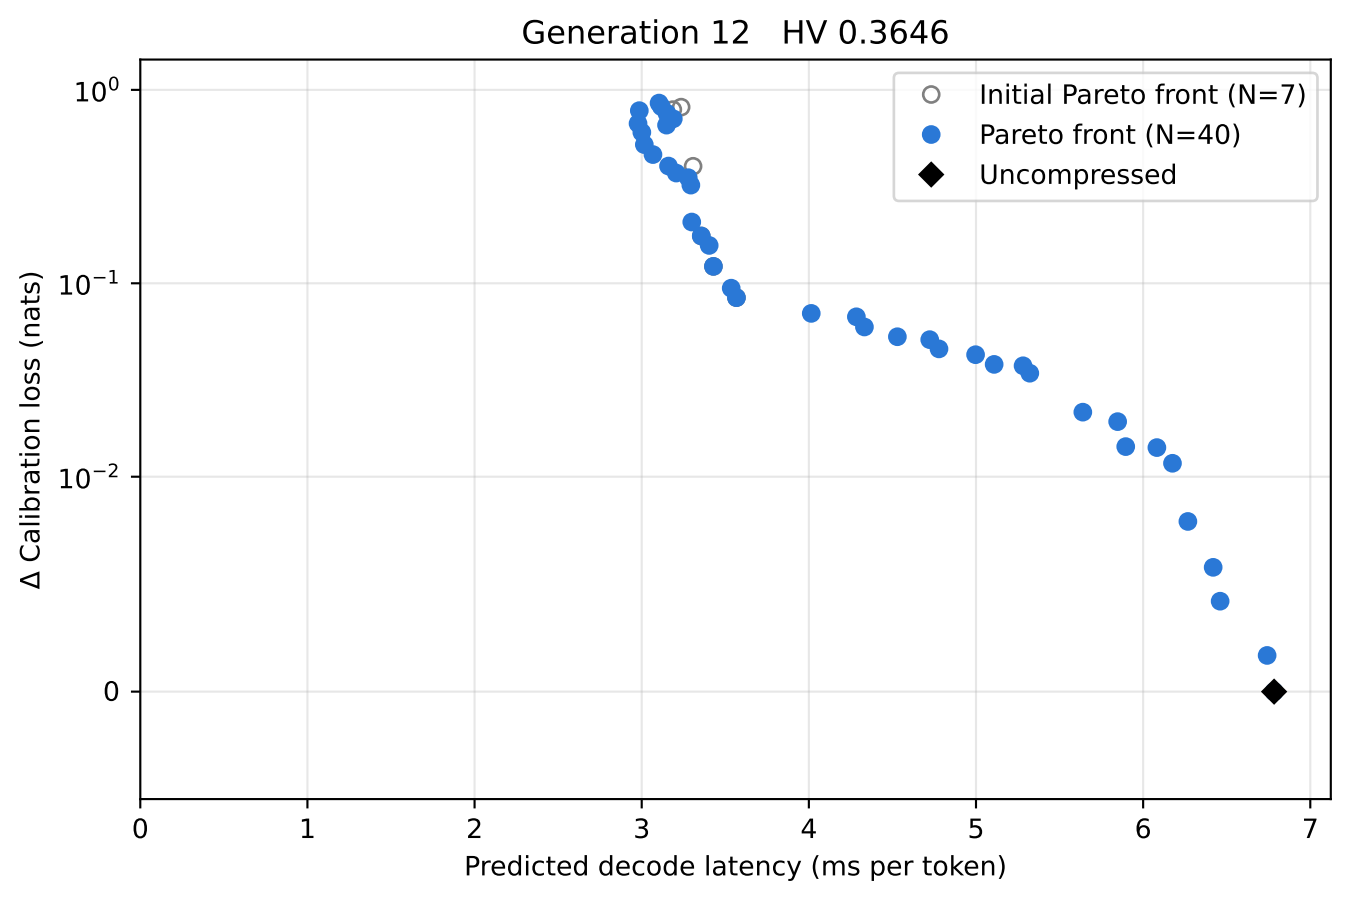
Generation 12 HV 0.3646
Initial Pareto front (N=7)
Pareto front (N=40)
Uncompressed
Δ Calibration loss (nats)
10^{0}
10^{-1}
10^{-2}
0
0
1
2
3
4
5
6
7
Predicted decode latency (ms per token)