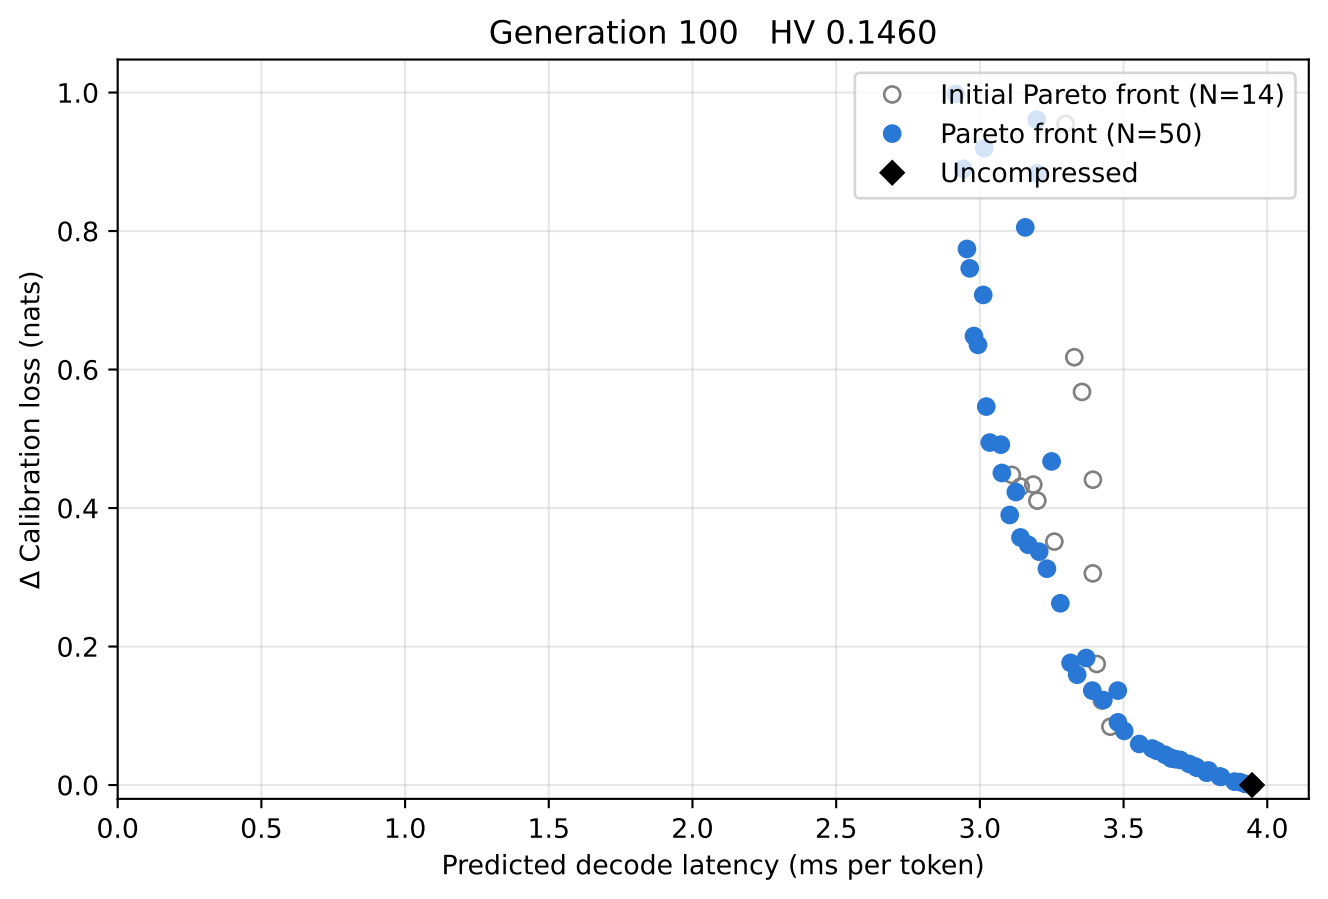
Generation 100 HV 0.1460
Initial Pareto front (N=14)
Pareto front (N=50)
Uncompressed
Δ Calibration loss (nats)
Predicted decode latency (ms per token)
0.0
0.2
0.4
0.6
0.8
1.0
0.0
0.5
1.0
1.5
2.0
2.5
3.0
3.5
4.0