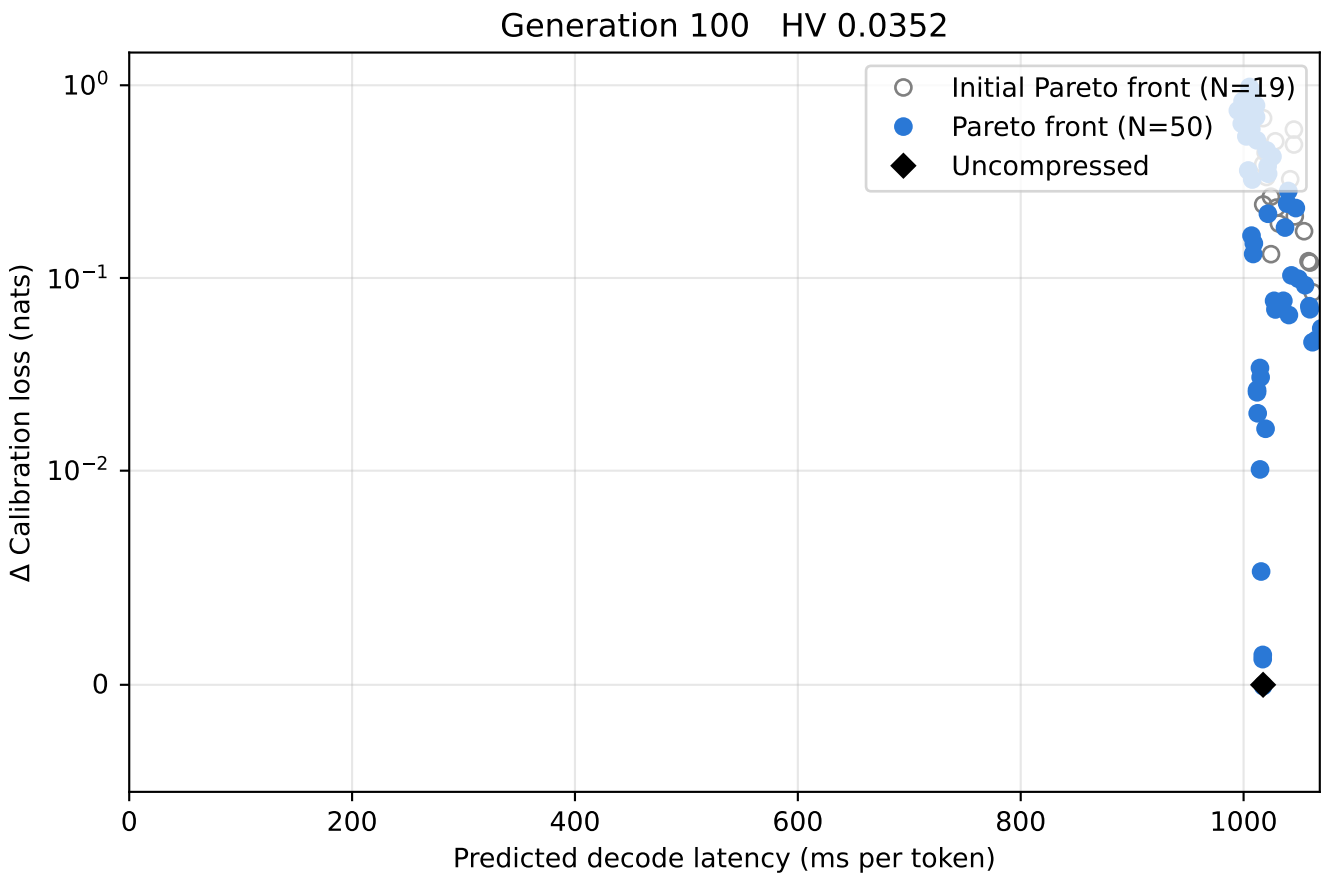
Generation 100 HV 0.0352
Initial Pareto front (N=19)
Pareto front (N=50)
Uncompressed
Δ Calibration loss (nats)
10^0
10^{-1}
10^{-2}
0
0
200
400
600
800
1000
Predicted decode latency (ms per token)