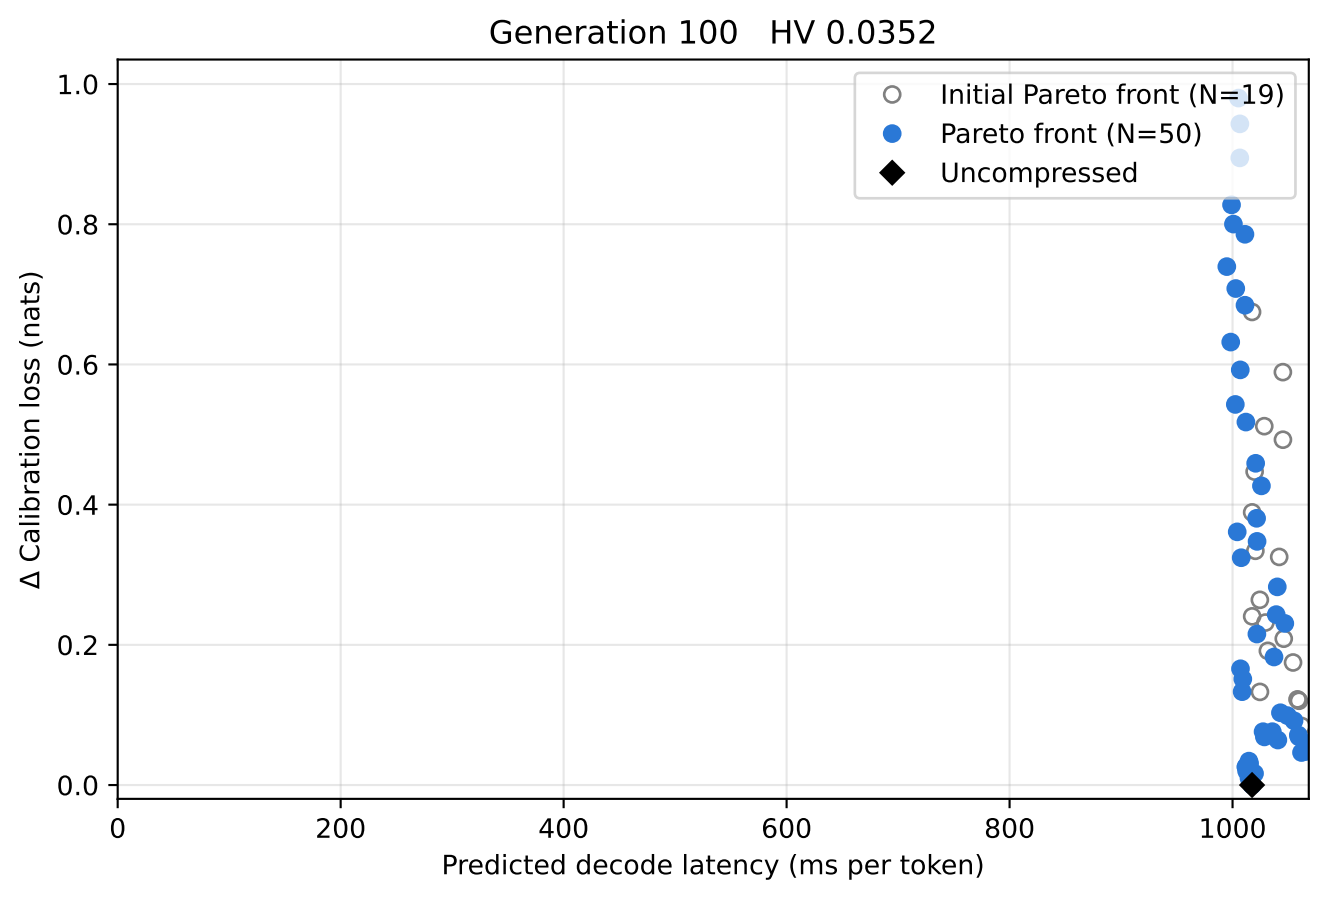
Generation 100 HV 0.0352
Initial Pareto front (N=19)
Pareto front (N=50)
Uncompressed
Δ Calibration loss (nats)
Predicted decode latency (ms per token)
0.0
0.2
0.4
0.6
0.8
1.0
0
200
400
600
800
1000