Hardware-Aware Multi-Objective Compression of Language Models: Minimizing Loss, Decode Latency and Memory Together

Rasa Khosrowshahli

November 2026

Abstract

Our earlier evolutionary compression search traded the loss of a language model against a storage size computed from idealized encodings, and it evaluated accuracy on a fake-quantized model whose weights are written back as dense bf16 tensors. Neither quantity changes how fast the model runs: a fake-quantized matrix takes exactly as long to multiply as the original, and unstructured sparsity at the levels the search produces is slower than dense on a GPU. This report describes a redesigned, hardware-aware search in which every candidate is a model the target GPU can execute faster. The decision variables are the weight format (bf16, or int4 with groups of 32, 64 or 128 weights) of each matrix role in each depth band, the format of the output head, the fraction of MLP neurons kept per band, and the number of attention and MLP sublayers removed. Parameters that change accuracy but not cost, such as the clipping range of each int4 group, are solved per group instead of being searched. The latency objective is the GPU time of one decode step, predicted from a table of kernel times measured on the target GPU; the memory objective is the deployed footprint of weights, scales and KV cache. On Qwen2.5-0.5B-Instruct and an NVIDIA RTX 3060, the latency model predicts the profiled kernel time of the deployed models within $\pm 1.1\%$, and the simulated loss matches the loss with the real int4 kernel within 0.0024 nats. Uniform int4 halves the decode time ($6.78 \rightarrow 3.4$–3.6 ms) at a loss increase of 0.08–0.18 nats; further speed comes only from structural pruning, at a much larger loss. A two-objective search with objectives {loss, latency} (memory as a constraint) and a three-objective search with objectives {loss, latency, memory} and the same budget produce nearly the same front (three-objective hypervolume 0.3655 vs 0.3646), because latency and memory are almost collinear on this model; the two-objective search converged faster in every objective space. Longer searches on an NVIDIA A100 80 GB (population 50, 100 generations) confirm the ordering of the two scenarios and show that the gains depend on the GPU: uniform int4 is only $1.16\times$ faster there, and the fastest solutions are $1.33\times$ faster in GPU time but only $1.05\times$ in eager wall-clock time. A third search mixes weight sharing and pruning with these formats and MLP width pruning, for the GPU time of an agent's turn at 32,000 tokens of context. Attention over the KV cache takes 66% of that turn, so its time moves by at most 2.2% below the uncompressed model. Sharing nearly halves the stored size at the lowest loss but barely changes GPU memory.

Contents

1 Introduction **3**

2 Motivation: what the GPU can accelerate **4**

3 Problem formulation **5**

3.1 Model structure 5

3.2 Decision variables . . . 5
3.3 What is solved instead of searched . . . 6
3.4 Objectives and constraints . . . 6

4 Compression operators **6**
4.1 int4 group quantization . . . 6
4.2 MLP width pruning . . . 7
4.3 Sublayer removal . . . 7
4.4 The simulated model is the deployed model . . . 7

5 Cost models **7**
5.1 Memory . . . 7
5.2 Decode latency . . . 8
5.3 Anatomy of a decode step . . . 8
5.4 Validation against the deployed models . . . 9
5.5 Measurement pitfalls . . . 9

6 Search algorithm **10**

7 Experiments on Qwen2.5-0.5B-Instruct **11**
7.1 Setup . . . 11
7.2 The fronts . . . 11
7.3 Test perplexity . . . 11
7.4 Two objectives against three . . . 13

8 Hardware-aware search with three objectives on an A100 **15**
8.1 Setup . . . 16
8.2 Anatomy of a decode step on the A100 . . . 16
8.3 The front . . . 16
8.4 Test perplexity . . . 16
8.5 Validation against the deployed models . . . 18

9 Hardware-aware search with two objectives on an A100 **18**
9.1 Setup . . . 18
9.2 The front . . . 18
9.3 Validation against the deployed models . . . 18
9.4 Two objectives against three . . . 20

10 Mixing weight sharing, pruning and architecture optimization **20**
10.1 Method . . . 20
10.2 The turn time barely changes . . . 23
10.3 Stored size against GPU memory . . . 23
10.4 What the search chose . . . 24
10.5 Test perplexity . . . 24
10.6 Search behavior . . . 26
10.7 Limitations . . . 26

11 Discussion **26**

12 Future work **27**

A Measured matrix-product times for Qwen2.5-0.5B **28**

B Sublayer removal ranking for Qwen2.5-0.5B **28**

1 Introduction

Post-training compression of language models is usually judged by the size of the compressed weights and by perplexity. What a deployment pays for, however, is the time to generate each token and the GPU memory the model occupies. Both fall only when the model is stored and executed in formats for which fast kernels exist, so a compression search that targets them has to know which formats the target GPU can execute and how fast.

Our previous search applied uniform binning quantization (UBQ), optionally followed by entropy-constrained merging of the levels (UBQ-ECSQ), with four or five genes shared by every matrix of the model: the number of grid levels k, a clipping factor c, a pruning interval $[\alpha, \beta]$ in units of the row's standard deviation, and the merge ratio ρ. Its second objective was the encoded size of the weights in the smallest of several storage formats (dense codes, a bitmap of non-zeros, CSR, or entropy-coded codes). Section 2 shows with measurements why this formulation cannot trade accuracy for speed: the compressed weights are evaluated as dense bf16 tensors, which run at the original speed; the storage formats have no fast decoding kernels; and unstructured pruning at the sparsity the search produces is slower than a dense matrix product.

This report describes a hardware-aware redesign built on four principles:

1. **Every candidate is executable.** The search space contains only formats with a kernel on the target GPU (a "kernel menu"); every configuration can be built and timed.
2. **Genes are the choices that change cost.** Parameters that change only accuracy at a fixed cost, such as the clipping range of a quantization group, have a single best value per layer and are solved deterministically inside the evaluation.
3. **Latency comes from a measured table, not from a stopwatch in the loop.** Kernel times are measured once per GPU and summed per configuration, which makes the objective deterministic and cheap.
4. **Decisions are per layer group,** since layers differ both in sensitivity and in how much time their matrices take.

The contributions are:

- a decision space of 23 integer genes for Qwen2.5-0.5B (24 for nanochat d34) mixing int4 quantization, MLP width pruning and sublayer removal, with per-group clipping solved by an inner solver (Sections 3 and 4);
- an additive decode-latency model built from CUDA-graph kernel timings and profiler-calibrated overheads, validated against the deployed models within ±1.1%, and a memory model of the deployed footprint (Section 5);
- an NSGA-II search with two variants, objectives {loss, latency} under a memory constraint and objectives {loss, latency, memory} (Section 6);
- a study on Qwen2.5-0.5B-Instruct on an RTX 3060 comparing the two variants by hypervolume (Section 7), repeated with longer searches on an A100 (Sections 8 and 9);

- a search that mixes weight sharing, pruning and hardware-aware formats, with the GPU time of an agent's turn as its objective (Section 10).

2 Motivation: what the GPU can accelerate

All measurements up to Section 7 were made on an NVIDIA GeForce RTX 3060 8 GB (Ampere, compute capability 8.6, 28 SMs, 1.5 MB L2 cache, 128-bit GDDR6 at 15 Gb/s, i.e. 240 GB/s), with PyTorch 2.13 and CUDA 12.6 on Windows 11, using only kernels that ship with PyTorch (no Triton, torchao, vLLM or FLUTE). Sections 8–10 use an NVIDIA A100 80 GB.

Quantization does not change latency. The search wrote the dequantized weights back into the dense bf16 matrices, and the forward pass multiplies them with a dense cuBLAS kernel, which reads every 16-bit weight whatever its value. Table 1 confirms that a matrix with 3 distinct values and 93% zeros is multiplied exactly as fast as the original.

Weights of an 8704×2176 matrix (d34 MLP)	Zeros	Batch 1 (μs)	Batch 2048 (μs)
Original bf16	0%	175.6	3092
UBQ, $k = 256$	1%	171.3	3053
UBQ, $k = 16$	21%	175.9	3055
UBQ, $k = 3$ with interval pruning	93%	176–181	3059–3140

Table 1: Time of one bf16 matrix product with fake-quantized weights (mean over 500 calls). The values of the weights do not affect a dense kernel.

The kernel menu. Table 2 times the formats that could reduce latency on the four matrix shapes of nanochat d34 (attention 2176×2176, MLP 8704×2176 and 2176×8704, output head 65536×2176) at a decode batch of 1 and 16 tokens and at a 2048-token prompt.

Format	Batch 1	Batch 16	2048 tokens	On this machine
bf16 with 25% of the rows removed (structured)	1.24–1.35×	1.29–1.78×	1.28–1.64×	yes, dense kernel
int4 weight-only, group 128 (tinygemm)	2.91–3.76×	1.51–1.83×	0.19–0.32×	yes, built into PyTorch
int8 weights and activations (eager)	0.20–1.60×	0.27–1.48×	0.50–0.77×	matmul only
int8 matrix product alone	–	–	1.78–2.86×	needs a fused kernel
bf16 2:4 semi-structured sparsity	–	–	–	no
CSR, 50% unstructured zeros	0.22–0.42×	≤0.02×	–	yes
CSR, 90% unstructured zeros	0.59–2.01×	0.09–0.27×	–	yes

Table 2: Speedup over a bf16 matrix product, as the range over the four d34 matrix shapes. Values below 1 are slowdowns. The int8 row with activations quantizes them in eager PyTorch, outside the kernel; the row below it times the int8 tensor-core product alone. 2:4 sparsity needs CUTLASS sparse kernels or cuSPARSELt, which the Windows build of PyTorch does not include.

Summed over d34's 34 blocks and output head, the matrix products of one decode step take 20.1 ms in bf16 (the dense kernels reach about 220 GB/s, 92% of the peak bandwidth), 7.1 ms with int4

blocks and a bf16 head, and 6.1 ms with int4 everywhere (3.3×). Four consequences shape the design:

- **Speedup depends on the workload.** Weight-only int4 is a large win when decoding one token at a time, which is limited by the bytes read from memory, and a 3–5× loss on long prompts, which are limited by compute. A latency objective is therefore defined for a workload.
- **Unstructured pruning does not buy speed** at the sparsity the earlier search reached; only structured pruning (rows, neurons, whole sublayers) or 2:4 patterns do.
- **Kernels execute fixed code widths.** A grid of $k = 37$ levels needs 6-bit codes and uses 37 of their 64 values; deployable formats use all 2^b levels of a supported width b.
- **Latency and memory are correlated when decoding,** because time is roughly the bytes read divided by the bandwidth. They diverge for memory that is not read per token (an embedding table, of which a lookup reads one row) and for the KV cache.

3 Problem formulation

3.1 Model structure

The search applies to HuggingFace Llama/Qwen2-style models and to nanochat GPT models [10]. A model has L blocks, each with an attention sublayer (projections q, k, v and the output projection o) and an MLP sublayer (an up projection, which is the pair gate and up in a SwiGLU MLP, and a down projection), followed by an output head. Matrices are grouped by *role* $r \in \{\mathrm{qkv}, \mathrm{o}, \mathrm{up}, \mathrm{down}\}$ and *depth band* $b \in \{1, \ldots, B\}$, the layers split into B equal runs. The q, k and v projections share a format because deployments fuse them into one product, and so do gate and up.

Qwen2.5-0.5B-Instruct [18] has $L = 24$ blocks of width $d = 896$, 14 query heads and 2 key-value heads of dimension 64 (grouped-query attention), a SwiGLU MLP of width $I = 4864$, and a vocabulary of 151,936 tokens whose embedding is tied to the output head (136 M of its 494 M parameters). All 493,961,216 matrix weights are compressible; the remaining 43,904 parameters are normalization weights.

3.2 Decision variables

Table 3 lists the genes. All are integers that index an ordered list of options, so that a step of one changes the option to its neighbor.

Gene	Granularity	Options	Count	What it changes
$f_{r,b}$	role × band (4×4)	bf16, int4 group 32, 64, 128	16	bytes read per token, memory
f_{head}	output head (with a tied embedding)	as above	1	the largest single product
f_{emb}	untied embeddings (d34 only)	as above	0 (1)	memory only
κ_b	MLP width per band	$1, \frac{7}{8}, \frac{3}{4}, \frac{5}{8}, \frac{1}{2}$	4	every regime
s_{attn}	attention sublayers removed	$0, \ldots, L/4$	1	products, attention, KV cache
s_{mlp}	MLP sublayers removed	$0, \ldots, L/4$	1	products

Table 3: Genes of the hardware-aware search: 23 for Qwen2.5-0.5B, 24 for nanochat d34.

The kept MLP width of a band is $n_b = 128 \cdot \mathrm{round}(\kappa_b I/128)$ neurons, a multiple of 128 so that the pruned matrices keep tensor-core-friendly shapes and whole int4 groups. With $B = 4$ bands and at most $L/4 = 6$ removals, the Qwen search space has $4^{17} \cdot 5^4 \cdot 7^2 \approx 5.3 \times 10^{14}$ configurations; for d34 ($L = 34$, an untied embedding) it has $4^{18} \cdot 5^4 \cdot 9^2 \approx 3.5 \times 10^{15}$.

3.3 What is solved instead of searched

Three choices change accuracy but not cost, so they are made by deterministic procedures computed once or cached per matrix (Section 4):

- the clipping range of each int4 group, chosen from six ratios by an activation-weighted error;
- which MLP neurons are kept, ranked by a Wanda-style importance score [20];
- which sublayers are removed first, ranked once by the loss increase of removing each alone.

3.4 Objectives and constraints

Let x be a gene vector and θ_x the model it describes. The accuracy objective is the increase of the mean next-token cross-entropy on a fixed calibration batch D,

$$\Delta\mathcal{L}(x) = \mathcal{L}_D(\theta_x) - \mathcal{L}_D(\theta_0), \qquad \mathcal{L}_D(\theta) = -\frac{1}{|D|}\sum_{t \in D} \log p_\theta(x_{t+1} \mid x_{\le t}), \tag{1}$$

measured in nats per token; a difference of $\Delta\mathcal{L}$ nats is a perplexity ratio of $e^{\Delta\mathcal{L}}$. With the decode latency $T(x)$ and the deployed memory $M(x)$ of Section 5, the two scenarios are

$$\text{2 objectives:} \quad \min_x \ (\Delta\mathcal{L}(x),\, T(x)) \qquad \text{s.t. } \Delta\mathcal{L}(x) \le 1,\ M(x) \le M_{\max}, \tag{2}$$

$$\text{3 objectives:} \quad \min_x \ (\Delta\mathcal{L}(x),\, T(x),\, M(x)) \qquad \text{s.t. } \Delta\mathcal{L}(x) \le 1,\ M(x) \le M_{\max}. \tag{3}$$

With a single target GPU, memory matters first as a constraint: a model that fits gets no faster by shrinking further. As an objective it adds the trade-off between memory and latency where the two diverge. Both scenarios log $M(x)$ for every solution, so their fronts can be compared in all three objectives.

4 Compression operators

4.1 int4 group quantization

The int4 format is the one consumed by PyTorch's built-in tinygemm kernel (compute capability 8.0 and above; also used by torchao [17]). A matrix $W \in \mathbb{R}^{N \times K}$ is split into groups of g consecutive weights of a row. For a group G with range $[w_{\min}, w_{\max}]$ and a clipping ratio γ,

$$s_\gamma = \mathrm{bf16}\Big(\frac{\gamma\,(w_{\max} - w_{\min})}{15}\Big), \qquad z_\gamma = \mathrm{bf16}\,(\gamma\, w_{\min} + 8\, s_\gamma), \tag{4}$$

$$q_j = \mathrm{clip}\Big(\,\mathrm{round}\,\Big(\frac{w_j - z_\gamma}{s_\gamma}\Big) + 8,\ 0,\ 15\Big), \qquad \hat{w}_j = (q_j - 8)\, s_\gamma + z_\gamma, \tag{5}$$

so that $q = 0$ and $q = 15$ reconstruct $\gamma w_{\min}$ and $\gamma w_{\max}$. The ratio of each group is

$$\gamma^\star = \arg \min_{\gamma \in \{1, 0.95, 0.9, 0.85, 0.8, 0.75\}} \sum_{j \in G} \mathbb{E}[x_j^2]\, \big(w_j - \hat{w}_j(\gamma)\big)^2, \tag{6}$$

where $\mathbb{E}[x_j^2]$ is the second moment of input channel j on the calibration batch. This is the output error $\mathbb{E}\|(W - \hat{W})x\|^2$ under a diagonal approximation of $\mathbb{E}[xx^\top]$, in the spirit of the clipping search of AWQ [12]. Error-compensating quantizers such as GPTQ [7] fit the same slot and are left for later. A group of g weights costs $g/2$ bytes of codes plus 4 bytes of scale and zero, i.e. $4 + 32/g$ bits per weight.

4.2 MLP width pruning

Neuron j of block ℓ is scored by $\|W^{\text{down}}_{\ell,:,j}\|_2 \sqrt{\mathbb{E}[h^2_{\ell,j}]}$, the norm of its output weights times the root mean square of its activation h (the input of the down projection), as in Wanda [20] aggregated per neuron (see also LLM-Pruner [13]). The n_b highest-scoring neurons are kept: the deployed model slices the rows of gate and up and the columns of down. In the search, the removed rows and columns are set to zero instead. For the down projection, the int4 groups are formed over the kept columns only, exactly as in the sliced matrix.

4.3 Sublayer removal

Removing a sublayer means that it adds nothing to the residual stream; in the search this is done by zeroing its output projection (o or down), which none of the supported models follows with a bias. The order of removal is fixed once: each of the $2L$ sublayers is removed alone and the loss increase is measured on 4 calibration windows, then s_{attn} removes the first s_{attn} attention sublayers of this ranking and s_{mlp} the first MLP sublayers, as in ShortGPT [14] with a direct loss criterion. Appendix B lists the ranking for Qwen2.5-0.5B. Reducing L binary genes to one integer per kind keeps the genome small.

4.4 The simulated model is the deployed model

The search evaluates the loss with the deployed weights written into the model: dequantized int4 values, zeroed neurons and zeroed output projections. The deployed model is built separately: int4 matrices on the tinygemm kernel, MLPs sliced to the kept neurons and removed sublayers replaced by a module that returns zeros. On small random Qwen2 and nanochat models the two agree exactly for pruning and removal in bf16, and to bf16 rounding (relative error below 3×10^{-2} on the logits) with the int4 kernel. On the real model the calibration loss of the simulated and deployed models differ by at most 0.0024 nats on all validated points (Table 5).

5 Cost models

5.1 Memory

A matrix of $N \times K$ weights takes $B(N, K, \text{bf16}) = 2NK$ bytes, or $B(N, K, \text{int4}_g) = NK/2 + 4N\,K/g$ bytes. The deployed memory of configuration x for β sequences of C tokens of context is

$$M(x) = 2P_{\text{rest}} + \sum_{m \in \mathcal{M}(x)} \Big(B(N_m, K_m, f_m) + 2N_m \mathbb{1}[\text{bias}_m] \Big) + \beta\, C\, (L - |A(x)|) \cdot 2H_{kv} d_h \cdot 2, \quad (7)$$

where $\mathcal{M}(x)$ are the matrices that remain (pruned to their kept shape), P_{rest} the remaining parameters in bf16, $A(x)$ the removed attention sublayers, and the last term the bf16 keys and values of the attention sublayers that remain. A tied embedding is counted once with the output head. For the uncompressed Qwen model at $\beta = 1$, $C = 1024$ this gives 988.1 MB of weights and 12.6 MB of KV cache.

5.2 Decode latency

The latency objective is the GPU time of one decode step for β sequences at C tokens of context, as a runtime without launch overhead (CUDA graphs) would see it:

$$T(x) = T_{\text{fixed}} + \sum_{\ell \notin A(x)} \Big(T^{\text{attn}}_{\text{misc}} + T_{\text{attn}}(\beta, C) + \sum_{m \in \{q,k,v,o\}} t(N_m, K_m, f_m) \Big) + \sum_{\ell \notin U(x)} \Big(T^{\text{mlp}}_{\text{misc}} + \sum_{m \in \text{up}_\ell} t(n_{b(\ell)}, d, f_m) + t(d, n_{b(\ell)}, f_{\text{down}}) \Big) + t(V, d, f_{\text{head}}), \tag{8}$$

where $U(x)$ are the removed MLP sublayers and $t(N, K, f)$ the measured time of a product of β tokens with an $N \times K$ matrix in format f.

Matrix products. t is measured once for every shape the deployed model can contain (all kept MLP widths included) and every format. Each measurement captures n calls in a CUDA graph, cycling through c copies of the weight with $c = \min(16, \lceil 4\,\text{L2}/\text{bytes} \rceil)$ so that the copies do not fit in the L2 cache together, as in the model, where every layer reads its own weights from memory. n is chosen so that one replay takes about 2 ms, which keeps the GPU at full clocks; the time is the minimum over 15 replays, since other work on the GPU can only slow a replay down. The GPU is warmed up for one second before the table is measured. A format whose kernel cannot run a shape falls back to bf16 for that shape.

Attention. T_{attn} is the faster of PyTorch's scaled-dot-product attention and a plain grouped-query implementation (two batched matrix products in which each query head reads its key-value head in place). Builds without FlashAttention [5], such as this Windows build, run SDPA's math path for grouped-query attention, which expands the cache: at 1024 tokens of context for Qwen it takes 130.6 μs per layer against 19.9 μs for the grouped implementation.

Everything else. Norms, rotary embeddings, activations, residual additions, the embedding lookup and the logits are calibrated on the real model. The profiler's summed kernel durations of one forward pass of β tokens without a cache (the fastest of three, after a warm-up) are measured for the full model (T_{full}), with every attention sublayer removed ($T_{\text{no-attn}}$) and with every MLP sublayer removed ($T_{\text{no-mlp}}$). Then

$$T^{\text{attn}}_{\text{misc}} = \frac{T_{\text{full}} - T_{\text{no-attn}} - \sum_{\text{attn}} t^{\text{bf16}}}{L} - T^{\text{own}}_{\text{attn}}(\beta, 1), \tag{9}$$

$$T^{\text{mlp}}_{\text{misc}} = \frac{T_{\text{full}} - T_{\text{no-mlp}} - \sum_{\text{mlp}} t^{\text{bf16}}}{L}, \qquad T_{\text{fixed}} = T_{\text{no-attn}} + T_{\text{no-mlp}} - T_{\text{full}} - t^{\text{bf16}}_{\text{head}}, \tag{10}$$

each clamped at zero, where $T^{\text{own}}_{\text{attn}}(\beta, 1)$ is the model's own attention over one token. By construction, Eq. (8) then reproduces the profiled time of the uncompressed model. For Qwen: $T_{\text{full}} = 6772.6\,\mu$s, $T_{\text{no-attn}} = 5070.0\,\mu$s, $T_{\text{no-mlp}} = 3570.2\,\mu$s, hence $T^{\text{attn}}_{\text{misc}} = 23.8\,\mu$s, $T^{\text{mlp}}_{\text{misc}} = 6.0\,\mu$s and $T_{\text{fixed}} = 666.0\,\mu$s. Appendix A lists the measured product times. The approach is that of latency lookup tables in hardware-aware architecture search [4, 22] and mixed-precision quantization [21, 24].

5.3 Anatomy of a decode step

Table 4 decomposes the predicted decode step of Qwen2.5-0.5B. In bf16 the MLP products take 45% of the time and the output head 18%. With int4 everywhere the products shrink by 3–4×,

and attention, the per-sublayer operations and the fixed part become 55% of the step. They bound what weight compression alone can achieve: $6783/1859 \approx 3.6\times$.

	bf16		int4, group 128	
Part of the step	μs	share	μs	share
MLP products (gate, up, down) $\times 24$	3060	45%	859	26%
Attention products (q, k, v, o) $\times 24$	662	10%	314	9%
Attention over 1024 cached tokens $\times 24$	478	7%	478	14%
Output head (151936×896)	1202	18%	321	10%
Norms, rotary, activations, residuals $\times 24$	715	11%	715	21%
Fixed (embedding, final norm, logits)	666	10%	666	20%
Total	6783		3353	

Table 4: Predicted GPU time of one decode step of Qwen2.5-0.5B-Instruct (batch 1, 1024 tokens of context) on the RTX 3060.

5.4 Validation against the deployed models

After each search, six points of the final front and the uncompressed model are built as deployed models. One forward pass of a single token (the shape the overheads were calibrated on) is profiled and compared with Eq. (8) evaluated for that shape. Table 5 shows the results for both searches of Section 7. The prediction is within -0.7% to $+1.1\%$ of the measured GPU time, including configurations with pruned MLPs and up to 10 removed sublayers. The last two columns compare the calibration loss (on two windows) of the simulated and the deployed model.

The same table shows the wall-clock time of the forward pass in eager PyTorch: 23–35 ms whatever the configuration, because a HuggingFace forward of a 0.5 B model at batch 1 is limited by Python and kernel-launch overhead, not by the GPU. The speedups predicted here are realized by runtimes without that overhead (CUDA graphs, compiled execution, or a serving engine).

5.5 Measurement pitfalls

Four effects can make such a table wrong by up to $2\times$; the measurement handles each of them.

- **Attention fallback.** Without FlashAttention, SDPA's math path made Qwen's attention $6.6\times$ slower than necessary and would have made removing attention look far more valuable than on a reasonable runtime. The table takes the fastest available implementation.
- **Clock drift and contention.** An idle GPU runs its memory clock far below the maximum, and kernels of a few microseconds timed in short bursts varied by 20–40%; another process on the GPU slowed some measurements and not others. Warm-up, replays of about 2 ms and the minimum over replays brought run-to-run differences below 1%. The table must be measured on an otherwise idle GPU: the overhead terms are differences of profiles, and contention distorts them.
- **L2 reuse.** Timing one small matrix repeatedly measures the L2 cache, not memory; cycling through copies larger than the cache removes this.
- **Eager overhead.** Wall-clock timing of eager forwards measures the CPU, not the GPU (Table 5); the profiler's kernel durations measure the GPU.

Search	Configuration	GPU time (ms) predicted	GPU time (ms) measured	error	Eager (ms)	Loss (nats) simulated	Loss (nats) deployed
2 obj.	uncompressed	6.77	6.79	−0.2%	28.0	2.8926	2.8926
	mixed int4, MLP $\frac{7}{8}$ in 2 bands, −6 attn, −4 MLP	2.85	2.82	+1.0%	23.1	3.7105	3.7110
	mixed bf16/int4, MLP $\frac{7}{8}$ in 1 band, −6 attn	3.15	3.13	+0.6%	23.2	3.3420	3.3428
	int4 group 32 everywhere	3.56	3.52	+1.1%	29.7	2.9784	2.9781
	mixed bf16/int4 (8.3 bits per weight)	4.60	4.60	0.0%	28.2	2.9467	2.9456
	mostly bf16 (12.3 bits per weight)	5.72	5.72	−0.1%	26.7	2.9147	2.9137
3 obj.	uncompressed	6.77	6.80	−0.4%	34.6	2.8926	2.8926
	int4 group 128, MLP $\frac{1}{2}$ in 1 band, −5 attn	2.97	2.96	+0.3%	27.2	3.4952	3.4958
	mixed int4, MLP $\frac{1}{2}$ in 2 bands, −3 MLP	3.14	3.10	+1.1%	31.7	3.6098	3.6122
	int4, mostly group 128	3.39	3.36	+0.9%	33.2	3.0372	3.0386
	mixed bf16/int4 (8.3 bits per weight)	4.52	4.51	+0.1%	31.6	2.9489	2.9494
	mostly bf16 (12.8 bits per weight)	5.84	5.86	−0.4%	32.9	2.9091	2.9095

Table 5: Predicted against measured GPU time of the deployed models (one token, batch 1), eager wall-clock time, and calibration loss of the simulated against the deployed model.

6 Search algorithm

The search is NSGA-II [6], implemented with pymoo [2] (Figure 1).

- **Encoding and operators.** Genes are integers in $[0, n_i - 1]$. Simulated binary crossover (probability 0.9, $\eta = 3$) and polynomial mutation ($\eta = 3$, per-variable probability $1/23$) act on real values that are rounded back to integers; the small distribution index gives steps of more than one option.
- **Initial population.** The first four individuals use one format everywhere (bf16, which is the original model, and int4 with groups of 32, 64 and 128) with nothing pruned. The others are random, except that each structural gene (κ_b, s_{attn}, s_{mlp}) is zero with probability $1/2$, since random pruning mostly breaks the model.
- **Duplicates and constraints.** Offspring that decode to an existing configuration are regenerated. Constraints use NSGA-II's constrained dominance: feasible solutions first, infeasible ones by their violation.
- **Evaluation.** An evaluation writes the configuration into the model, measures the calibration loss, restores the original weights and evaluates $T(x)$ and $M(x)$ analytically.
- **Hypervolume.** The hypervolume [25] of each generation's front is normalized by the nadir of the first feasible front, with reference point $(1, \ldots, 1)$.

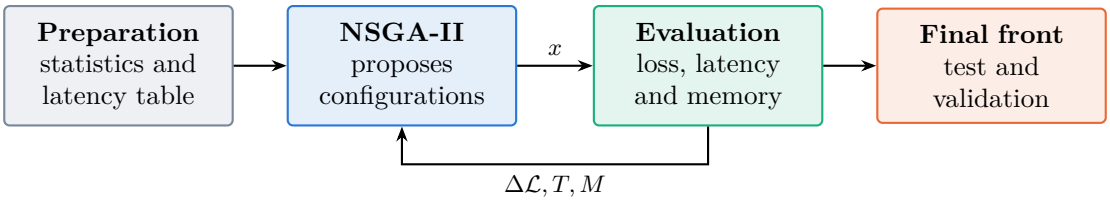


Figure 1: Pipeline of the hardware-aware search. Before the search, the activation statistics, the sublayer removal ranking and the latency table of the GPU are computed once. NSGA-II proposes configurations x; each is evaluated by its calibration loss $\Delta\mathcal{L}$ and its predicted latency T and memory M. The final front is scored on test sets and validated with deployed models.

7 Experiments on Qwen2.5-0.5B-Instruct

7.1 Setup

- **Calibration.** One fixed batch of 8 windows of 2048 tokens of C4 train [19]; the uncompressed model's loss on it is 2.8854 nats.
- **Workload.** Decoding one sequence at 1024 tokens of context; the uncompressed model's predicted step takes 6.78 ms and needs 1000.6 MB.
- **Constraints.** $\Delta\mathcal{L} \leq 1$ nat and $M \leq 7000$ MB. The memory constraint never binds: the uncompressed model needs 1001 MB.
- **Budget.** A population of 40 for 12 generations (480 evaluations) and seed 42 for both scenarios, which therefore start from the same initial population. An evaluation takes about 4.4 s, a generation about 175 s, a search about 36 minutes.
- **Preparation.** Activation statistics (4 s), the removal ranking (96 evaluations on 4 windows, 99 s) and the latency table (7 s) are computed once per run.
- **Test.** Six points of each final front, spread evenly by latency, are scored on 32 windows of 2048 tokens of WikiText-2 test [15] and of C4 validation.

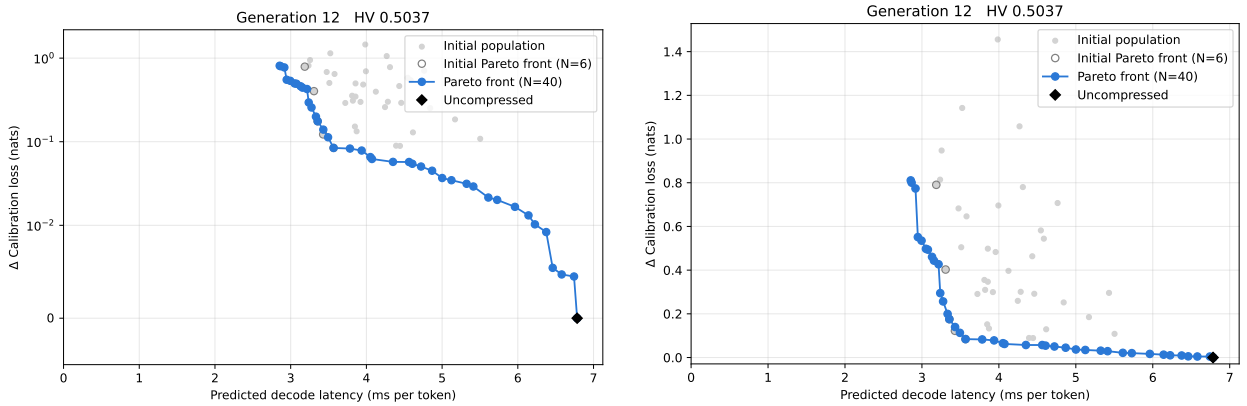
7.2 The fronts

Figure 2 shows how far the two-objective search moved from its initial population, and Figure 3 both final fronts in the latency and memory projections. The fronts have three regimes:

- From 6.8 ms down to about 3.6 ms, matrices move from bf16 to int4 at a small loss: int4 with groups of 32 everywhere costs 0.084 nats at 3.57 ms and 321 MB.
- Around 3.4–3.6 ms, coarser groups and the output head in int4 give the last cheap gains: group 128 everywhere costs 0.18 nats at 3.36 ms.
- Below 3.3 ms, only structural pruning remains: removing 6 attention and 4 MLP sublayers and pruning two MLP bands to $\frac{7}{8}$ reaches 2.85 ms at 0.81 nats. Without recovery fine-tuning, structural pruning appears only at this end of the front.

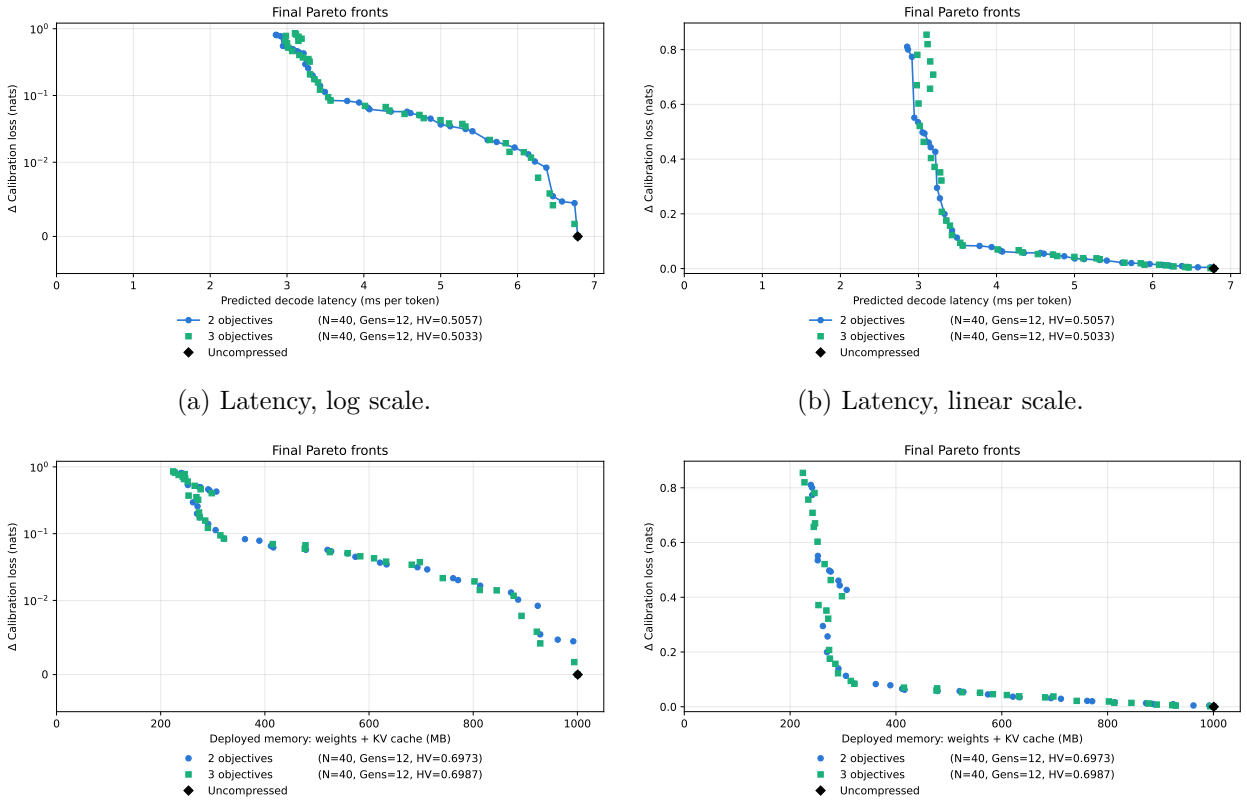
7.3 Test perplexity

Table 6 lists the test perplexity of the scored points, and Figure 4 shows WikiText-2 against latency. The ranking of the calibration loss carries over to both test sets. The knee of the front, int4



(a) Log scale. (b) Linear scale.

Figure 2: Two-objective search: initial population (grey), initial front (hollow) and final front after 12 generations, with the loss on a log scale (linear below 0.01 nats) and on a linear scale.



(a) Latency, log scale. (b) Latency, linear scale.

(c) Memory, log scale. (d) Memory, linear scale.

Figure 3: Final fronts of the two searches: loss against predicted decode latency (top) and against deployed memory (bottom), on a log scale (linear below 0.01 nats, left) and a linear scale (right). A front is drawn as a line only in the objectives it optimized; elsewhere its points are not ordered along the axis.

everywhere, costs 10% WikiText-2 and 9% C4 perplexity for a 1.9× faster decode step and 3.1× less memory.

Search	Latency (ms)	Memory (MB)	Bits/weight	$\Delta\mathcal{L}$ (nats)	WikiText-2	C4
Uncompressed	6.78	1001	16.00	0	14.33	22.33
2 obj.	5.73	771	12.28	0.020	14.67	22.80
	4.61	527	8.34	0.054	15.00	23.55
	3.57	321	5.00	0.084	15.73	24.29
	3.16	294	4.60	0.443	25.57	35.59
	2.85	239	3.72	0.811	39.12	51.80
3 obj.	5.85	802	12.79	0.019	14.56	22.83
	4.53	525	8.29	0.053	15.40	23.42
	3.40	286	4.42	0.157	17.18	26.78
	3.15	234	3.59	0.757	44.46	54.86
	2.98	247	3.84	0.670	32.87	49.46

Table 6: Test perplexity (32 windows of 2048 tokens each) of the uncompressed model and of points of the final fronts of the two searches.

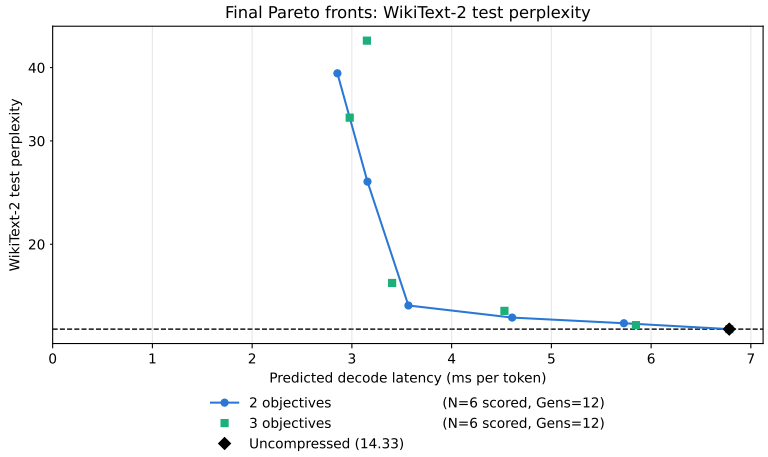


Figure 4: WikiText-2 test perplexity of the scored points against predicted decode latency.

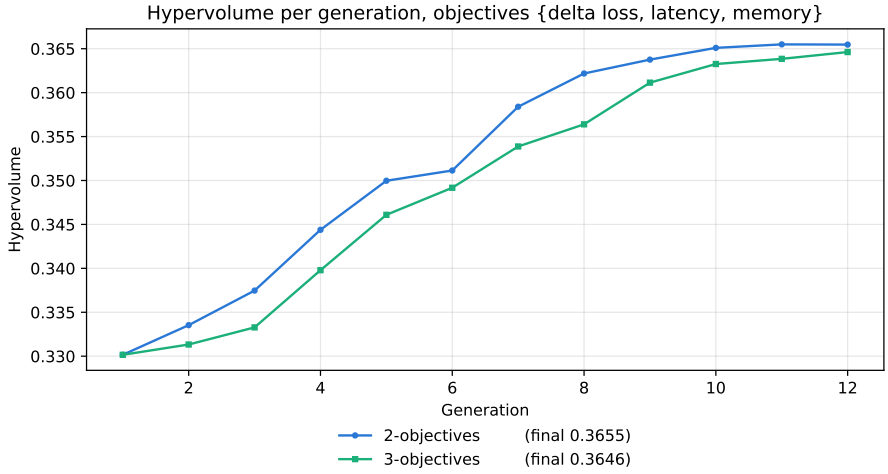
7.4 Two objectives against three

To compare the scenarios, both final fronts and every generation's front are evaluated in all three objectives with one normalization: each objective is divided by the worst value over the initial fronts of both runs, (0.814 nats, 6.78 ms, 1000.6 MB), with reference point $(1, 1, 1)$. The same is done in the two two-dimensional projections. Table 7 and Figure 5 show the results.
Three observations:

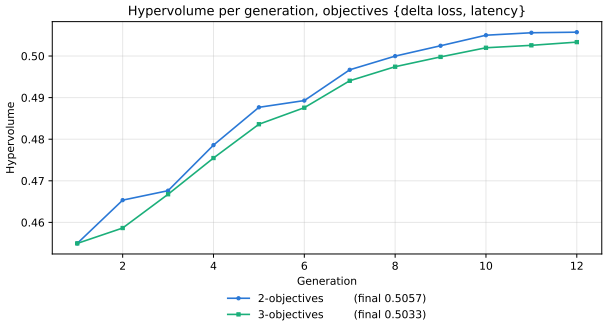
- **In all three objectives, the two-objective search is ahead in every generation**, by up to 1.6% at generation 8 and by 0.25% at the end. The ordering holds with a run-independent box (the $\Delta\mathcal{L}$ limit and the uncompressed latency and memory as the reference point: 0.3795 against 0.3776) and over every feasible solution each search evaluated, not only its final front (0.3815 with 76 non-dominated points against 0.3781 with 62).

Objectives	{loss, latency, memory}		{loss, latency}		{loss, memory}	
Generation	2 obj.	3 obj.	2 obj.	3 obj.	2 obj.	3 obj.
1	0.3302	0.3302	0.4550	0.4550	0.6504	0.6504
2	**0.3335**	0.3313	**0.4654**	0.4587	**0.6650**	0.6555
3	**0.3375**	0.3333	**0.4676**	0.4668	**0.6696**	0.6671
4	**0.3444**	0.3398	**0.4786**	0.4755	**0.6815**	0.6691
5	**0.3500**	0.3461	**0.4877**	0.4836	**0.6833**	0.6757
6	**0.3511**	0.3492	**0.4893**	0.4876	**0.6837**	0.6788
7	**0.3584**	0.3539	**0.4967**	0.4941	**0.6892**	0.6845
8	**0.3622**	0.3564	**0.5000**	0.4974	**0.6942**	0.6870
9	**0.3638**	0.3611	**0.5025**	0.4998	**0.6946**	0.6939
10	**0.3651**	0.3633	**0.5050**	0.5020	**0.6961**	0.6960
11	**0.3655**	0.3639	**0.5056**	0.5026	0.6971	**0.6974**
12	**0.3655**	0.3646	**0.5057**	0.5033	0.6973	**0.6987**

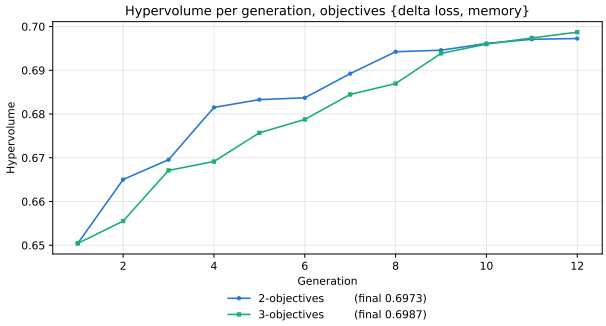
Table 7: Hypervolume of each generation's front with one shared normalization, in all three objectives and in the two projections. Generation 1 is identical because both searches start from the same population.



(a) {loss, latency, memory}.



(b) {loss, latency}.



(c) {loss, memory}.

Figure 5: Hypervolume per generation of the two searches with one shared normalization: in all three objectives (a), and in the projections {loss, latency} (b) and {loss, memory} (c).

- **Even in {loss, memory},** which only the three-objective search optimizes, the two-objective search leads until generation 10; the three-objective search overtakes it at generations 11–12, by 0.2%.
- **The fronts lie on a curve, not a surface** (Figure 6): memory and latency decrease together.

On this model the third objective is nearly redundant with the second, which is the situation analyzed by objective reduction [3]. A redundant objective adds no trade-off to find but spreads NSGA-II's diversity preservation over an extra dimension, which slows convergence along the front that matters. With one seed per scenario, differences of this size are within the usual seed-to-seed variation and are not significant.

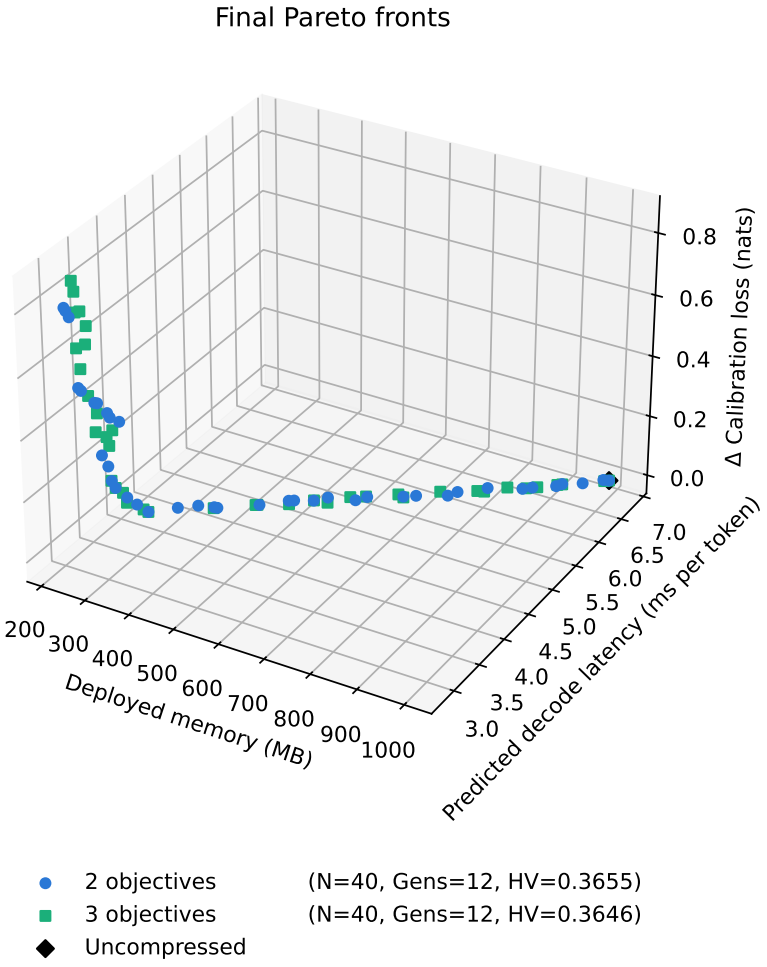


Figure 6: Final fronts of both searches in all three objectives. Both lie along a thin curve because memory and latency decrease together on this model.

8 Hardware-aware search with three objectives on an A100

This section and the next repeat the two scenarios of Section 7 with a longer search on a server GPU.

8.1 Setup

- **Hardware.** NVIDIA A100 80 GB PCIe, PyTorch 2.13 with CUDA 12.6 on Linux. The latency table is measured on this GPU; both searches use the same table.
- **Search.** Objectives {loss, latency, memory}, the same 23 genes, population 50, 100 generations, seed 42. The search evaluated 4953 distinct configurations, about 0.26 s each, in 21 minutes.
- **Calibration and constraints.** As in Section 7: the same 8 windows of 2048 tokens (uncompressed loss 2.8857 nats on this machine), $\Delta\mathcal{L} \leq 1$ nat, $M \leq 7000$ MB.
- **Workload.** Decoding one sequence at 1024 tokens of context. The uncompressed model's predicted step takes 3.95 ms and needs 1000.6 MB.
- **Test.** Six points of the final front on 32 windows of 2048 tokens of WikiText-2 test and of C4 validation, then six points built as deployed models.

8.2 Anatomy of a decode step on the A100

Table 8 repeats Table 4 on the A100. The step takes 3.95 ms against 6.78 ms on the RTX 3060, and the matrix products take 35% of it against 73%. Int4 everywhere therefore reaches only $3947/3407 = 1.16\times$, against about $2\times$ on the RTX 3060. Even free matrix products would reach only $3947/2546 = 1.55\times$.

	bf16		int4, group 128	
Part of the step	μs	share	μs	share
MLP products (gate, up, down) $\times 24$	719	18%	425	12%
Attention products (q, k, v, o) $\times 24$	520	13%	354	10%
Attention over 1024 cached tokens $\times 24$	301	8%	301	9%
Output head (151936×896)	163	4%	82	2%
Norms, rotary, activations, residuals $\times 24$	921	23%	921	27%
Fixed (embedding, final norm, logits)	1324	34%	1324	39%
Total	3947		3407	

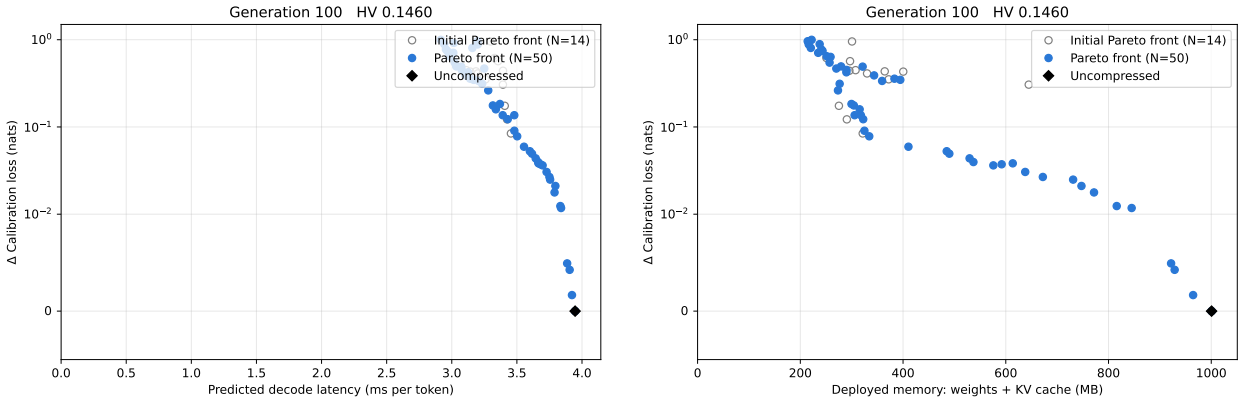
Table 8: Predicted GPU time of one decode step of Qwen2.5-0.5B-Instruct (batch 1, 1024 tokens of context) on the A100.

8.3 The front

The final front has 50 solutions from 2.92 to 3.95 ms and from 214 to 1001 MB (Figure 7). Formats alone give most of the memory and little of the time: int4 with groups of 32 everywhere costs 0.084 nats at 3.45 ms and 321 MB, and groups of 128 cost 0.175 nats at 3.41 ms and 275 MB. Going faster takes structural changes: of the 50 solutions, 28 remove attention sublayers, 14 remove MLP sublayers and 23 prune MLP neurons.

8.4 Test perplexity

Table 9 and Figure 8 give the scored points. The test sets follow the calibration loss. The point at 0.091 nats costs 11% WikiText-2 and 10% C4 perplexity for $3.1\times$ less memory and a $1.13\times$ faster decode step.



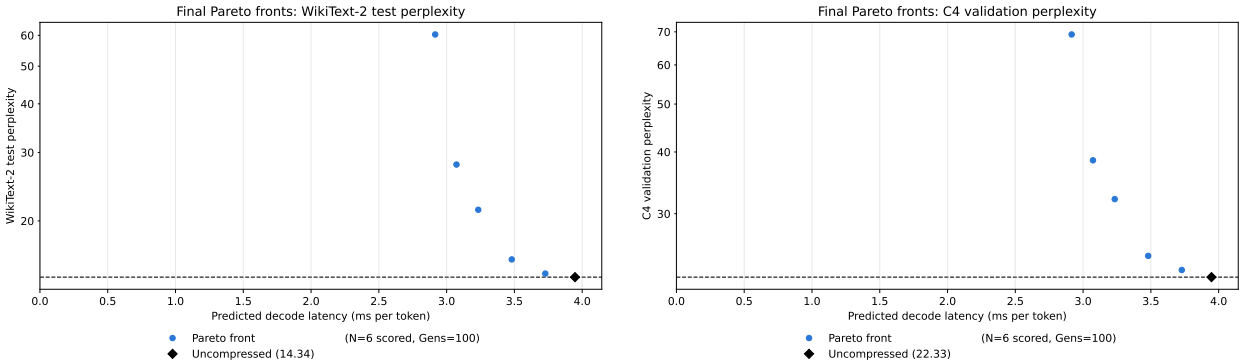
(a) Against predicted decode latency.

(b) Against deployed memory.

Figure 7: Three-objective search on the A100: final front (50 solutions), initial front and the uncompressed model.

Latency (ms)	Memory (MB)	$\Delta\mathcal{L}$ (nats)	WikiText-2	C4	Configuration
3.95	1001	0	14.34	22.33	uncompressed
3.73	638	0.030	14.64	23.08	int4 in up and part of down
3.48	325	0.091	15.92	24.65	mostly int4
3.23	277	0.312	21.37	32.11	int4, MLP pruned in 2 bands, −3 attn
3.07	321	0.492	27.93	38.47	mostly int4, MLP 5/8 in 1 band, −6 attn
2.92	222	0.998	60.31	69.15	int4, MLP pruned in 2 bands, −6 attn, −4 MLP

Table 9: Test perplexity (32 windows of 2048 tokens each) of the uncompressed model and of points of the final three-objective front.



(a) WikiText-2 test.

(b) C4 validation.

Figure 8: Test perplexity of the scored points of the three-objective search against predicted decode latency.

8.5 Validation against the deployed models

Table 10 compares six deployed models with the prediction, as in Table 5. The simulated and deployed losses differ by at most 0.003 nats. The latency model is 0.6–2.7% below the measured GPU time, a small systematic underestimate that did not appear on the RTX 3060.

The eager wall-clock column limits what these gains mean in practice. An eager forward takes 10.8–12.4 ms whatever the configuration, against 2.9–3.9 ms of GPU time, so a GPU-time gain of 1.32× is only 1.06× on the clock, and the compressed models without sublayer removal are slower than the uncompressed model (0.92–0.95×).

	GPU time (ms)			Eager	Loss (nats)		Speedup	
Configuration	pred.	meas.	error	(ms)	sim.	deployed	GPU	eager
int4, MLP pruned in 2 bands, −6 attn, −4 MLP	2.83	2.90	−2.7%	10.81	3.9098	3.9119	1.32	1.06
mostly int4, MLP 5/8 in 1 band, −6 attn	2.98	3.05	−2.1%	11.21	3.4023	3.4052	1.26	1.02
int4, MLP pruned in 2 bands, −3 attn	3.13	3.17	−1.4%	11.99	3.1785	3.1786	1.21	0.95
mostly int4	3.36	3.44	−2.3%	12.37	2.9839	2.9832	1.11	0.92
bf16 attention and head, int4 MLP	3.61	3.63	−0.6%	12.04	2.9206	2.9202	1.05	0.95
uncompressed	3.83	3.87	−1.1%	11.39	2.8916	2.8916	0.99	1.00

Table 10: Three-objective search: predicted against measured GPU time of the deployed models (one token, batch 1), eager wall-clock time, calibration loss on two windows, and the measured speedup over a separately timed uncompressed reference (3.83 ms, 11.41 ms eager).

9 Hardware-aware search with two objectives on an A100

9.1 Setup

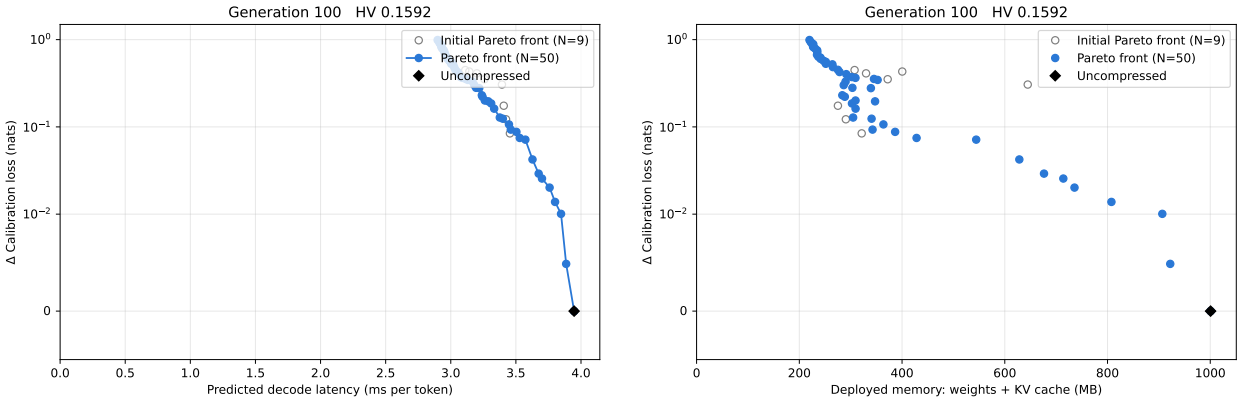
The hardware, calibration batch, constraints, workload, budget and test sets are those of Section 8; memory is only a constraint, which never binds. The search evaluated 4975 distinct configurations, about 0.27 s each, in 22 minutes.

9.2 The front

The final front has 50 solutions from 2.90 to 3.95 ms and from 220 to 1001 MB (Figure 9). Of the 50, 38 remove attention sublayers, 20 remove MLP sublayers and 23 prune MLP neurons. Table 11 gives the scored points. The test sets follow the calibration loss here too: the point at 0.088 nats costs 9% WikiText-2 and 8% C4 perplexity at 3.50 ms and 387 MB.

9.3 Validation against the deployed models

Table 12 shows the same pattern as Section 8.5: the latency model is 1.0–2.9% below the measured GPU time, the losses differ by at most 0.002 nats, and a GPU-time gain of up to 1.33× is at most 1.05× on the eager clock.

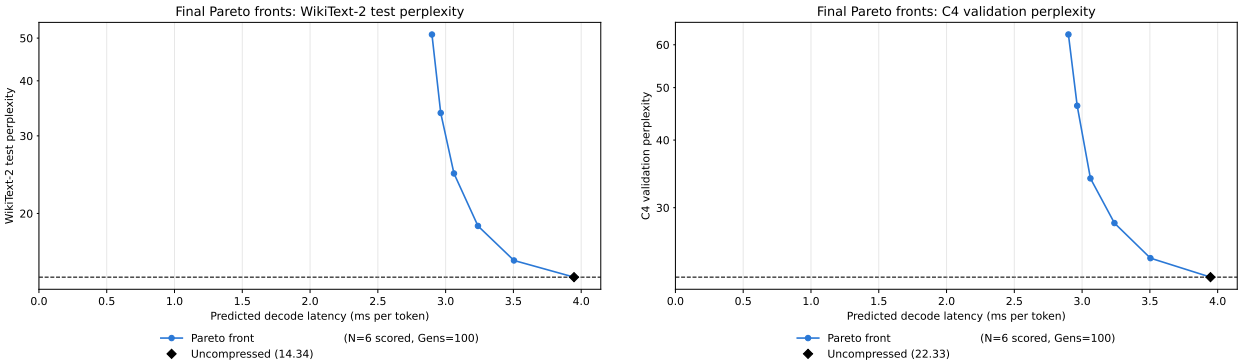


(a) Against predicted decode latency (objective). (b) Against deployed memory (reported).

Figure 9: Two-objective search on the A100: final front (50 solutions), initial front and the uncompressed model.

Latency (ms)	Memory (MB)	$\Delta\mathcal{L}$ (nats)	WikiText-2	C4	Configuration
3.95	1001	0	14.34	22.33	uncompressed
3.50	387	0.088	15.65	24.21	mostly int4
3.24	284	0.230	18.74	28.10	int4, −3 attn
3.06	291	0.402	24.65	33.99	int4, MLP 7/8 in 1 band, −6 attn
2.96	237	0.636	33.82	46.29	int4, MLP 5/8 in 2 bands, −6 attn, −1 MLP
2.90	220	0.990	50.95	62.69	int4, MLP 5/8 in 2 bands, −6 attn, −5 MLP

Table 11: Test perplexity (32 windows of 2048 tokens each) of the uncompressed model and of points of the final two-objective front.



(a) WikiText-2 test. (b) C4 validation.

Figure 10: Test perplexity of the scored points of the two-objective search against predicted decode latency.

	GPU time (ms)			Eager	Loss (nats)		Speedup	
Configuration	pred.	meas.	error	(ms)	sim.	deployed	GPU	eager
int4, MLP 5/8 in 2 bands, −6 attn, −5 MLP	2.81	2.89	−2.7%	11.22	3.8935	3.8947	1.33	1.05
int4, MLP 5/8 in 2 bands, −6 attn, −1 MLP	2.87	2.96	−2.9%	11.53	3.5279	3.5274	1.29	1.03
int4, MLP 7/8 in 1 band, −6 attn	2.97	3.00	−1.1%	11.65	3.3040	3.3042	1.28	1.02
int4, −3 attn	3.13	3.22	−2.7%	12.34	3.1046	3.1054	1.19	0.96
mostly int4	3.38	3.44	−1.6%	13.01	2.9748	2.9747	1.11	0.91
uncompressed	3.83	3.87	−1.0%	11.89	2.8916	2.8916	0.99	1.00

Table 12: Two-objective search: predicted against measured GPU time of the deployed models (one token, batch 1), eager wall-clock time, calibration loss on two windows, and the measured speedup over a separately timed uncompressed reference (3.83 ms, 11.83 ms eager).

9.4 Two objectives against three

Table 13 and Figure 11 compare both searches. Each objective is divided by one shared value (0.955 nats, 3.95 ms, 1000.6 MB), with reference point $(1, 1, 1)$.

- **In all three objectives, the two-objective search is ahead from generation 8 on,** and ends 3.7% higher (0.1514 against 0.1460). The lead shrinks to 1.5% over every feasible solution each search evaluated (0.1567 with 580 non-dominated points against 0.1543 with 426). The same ordering appeared on the RTX 3060 (Table 7).
- **In {loss, memory}, the three-objective search is ahead,** by 1.2% at the end.
- **The practical differences are small and uneven** (Table 14). Latency agrees within 1.2%. The three-objective search has less memory at low loss, up to 10% at $\Delta\mathcal{L} \leq 0.05$, and about the same elsewhere.
- **100 generations were more than needed.** In all three objectives, the two-objective search reaches its final level by generation 50, and its hypervolume drifts slightly down afterwards, since the front is limited to the current population.

Against the short runs of Section 7, the tenfold budget mainly improves the low-loss end: the smallest weights at $\Delta\mathcal{L} \leq 0.05$ fall from 561 to 476 MB (two objectives) and from 555 to 430 MB (three objectives), and change little at 0.1–0.3 nats. Times cannot be compared across the two GPUs. With one seed per scenario, these differences are indicative only.

10 Mixing weight sharing, pruning and architecture optimization

The searches above never touch the storage formats of Section 2, and they time one decode step. This search brings weight sharing and pruning back into one optimization with the deployed formats and the width of the MLPs, and it times the work of an agent: a turn with prompt tokens and generated tokens, over a long KV cache.

10.1 Method

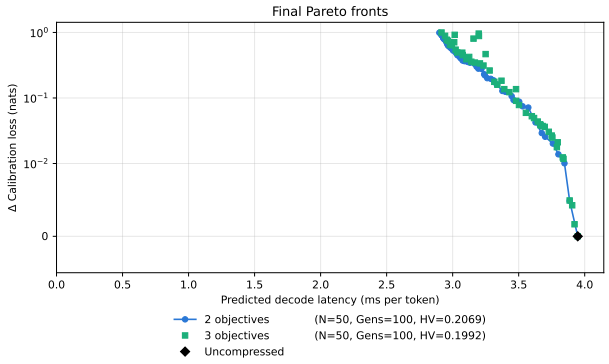
One format per matrix group. Stacking the methods does not work: measured on this model, codes of a shared grid pushed through the int4 grid lose their sharing and add loss, and pruning on

Objectives	{loss, latency, memory}		{loss, latency}		{loss, memory}	
Generation	2 obj.	3 obj.	2 obj.	3 obj.	2 obj.	3 obj.
1	0.1202	**0.1218**	0.1686	0.1686	0.6584	**0.6670**
5	0.1356	**0.1381**	0.1902	**0.1906**	0.6783	**0.6938**
10	**0.1455**	0.1403	**0.2009**	0.1926	0.6960	**0.7039**
25	**0.1497**	0.1406	**0.2051**	0.1946	**0.7004**	0.6971
50	**0.1516**	0.1422	**0.2063**	0.1941	0.7055	**0.7099**
75	**0.1515**	0.1439	**0.2068**	0.1958	0.7032	**0.7116**
100	**0.1514**	0.1460	**0.2069**	0.1992	0.6989	**0.7072**

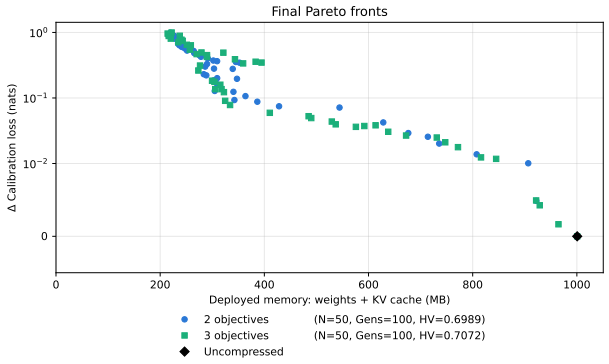
Table 13: Hypervolume of the fronts of the A100 searches with one shared normalization, in all three objectives and in the two projections.

	Fastest (ms)		Smallest weights (MB)	
$\Delta\mathcal{L}$ up to (nats)	2 obj.	3 obj.	2 obj.	3 obj.
0.02	3.76	3.77	756	725
0.05	3.59	3.60	476	430
0.1	3.44	3.44	295	299
0.2	3.26	3.30	263	263
0.3	3.18	3.19	250	249
0.5	3.02	3.03	229	228
1.0	2.90	2.92	207	202

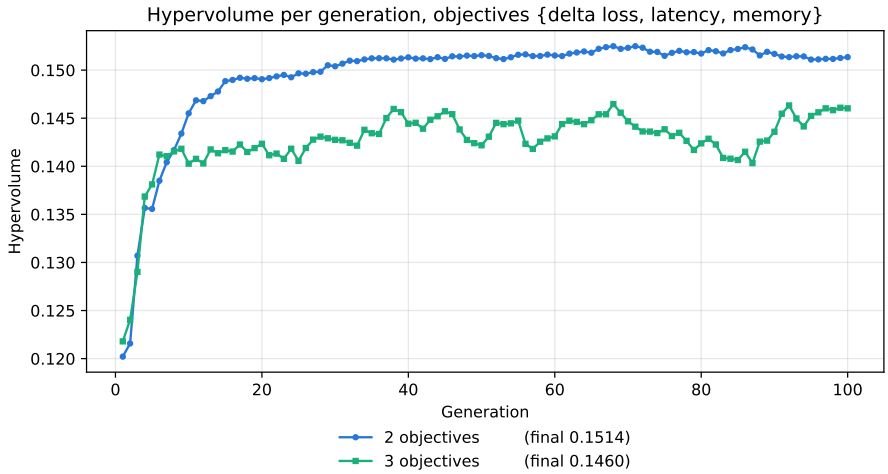
Table 14: Best latency and best deployed weights within each loss budget, over all evaluated solutions of each search. Each column is its own solution: the fastest and the smallest are usually different models.



(a) Latency.



(b) Memory.



(c) Hypervolume in {loss, latency, memory}.

Figure 11: Final fronts of the two A100 searches against predicted decode latency (a) and deployed memory (b), and their hypervolume per generation with one shared normalization (c).

top of int4 adds loss without shrinking the codes. So a group has one format: bf16, int4 with groups of 32, 64 or 128, or *shared*. A shared matrix is compressed by UBQ-ECSQ with one codebook for all shared matrices of the model. No GPU kernel reads these codes, so a shared matrix is expanded and runs as bf16: sharing changes the stored size and the accuracy, not the time or the GPU memory. Within a configuration the order is fixed: neurons are removed first, each kept matrix is put in its format, and the codes are entropy coded for storage last.

Genes. The 16 role and band genes and the head gene take one of the five formats. Four genes hold the MLP width per band, the only architecture decision in this run, since sublayer removal was turned off ($s_{\text{attn}} = s_{\text{mlp}} = 0$). Four genes hold the one sharing candidate: grid levels $k \in \{17, 33, 65, 129, 257\}$, clipping $c \in \{1, 0.9, 0.8, 0.7\}$, the pruning interval $\beta \in \{0, 0.1, 0.2, 0.3, 0.5\}$ in row standard deviations and the merge ratio $\rho \in \{1, 0.75, 0.5, 0.25\}$; they are zeroed when no group is shared. The search space has $5^{21} \cdot 400 \approx 1.9 \times 10^{17}$ configurations.

Objectives. $\Delta\mathcal{L}$, the stored size $S(x)$ and the GPU time of one agent's turn,

$$T_{\text{turn}}(x) = N_{\text{new}} \frac{T_{256}(x) - t_{\text{head}}(x)}{256} + N_{\text{gen}} \frac{T_4(x)}{4}, \tag{11}$$

where T_β is Eq. (8) at batch β from a table measured at that batch. Only decoding needs logits for every token, so the head is left out of the prompt term. Each decoding sequence reads its own cache of 32,000 tokens, while prompt tokens read one shared cache together in chunks of 256. The workload is $N_{\text{new}} = 2000$ new prompt tokens and $N_{\text{gen}} = 300$ generated tokens per turn, with 4 agents decoding together. $S(x)$ is the size under an ideal entropy coder, plus scales, tables and the codebook; it is a computed size, not a written file. Deployed memory, with the KV cache of the 4 agents, and the number of agents that fit on 80 GB with 10% kept free are reported, not optimized.

Search. $\Delta\mathcal{L} \leq 1$ nat, the same calibration batch, population 50, 100 generations, seed 42, on the A100: 4992 distinct configurations, about 0.32 s each, in 27 minutes.

10.2 The turn time barely changes

The uncompressed turn takes 1018 ms (303 μs per prompt token, 1369 μs per generated token), and the final front spans only 995–1070 ms (Figure 12). Table 15 shows why: attention over the KV cache is 66% of the turn and all weight products together are 13%, so weight compression has little to work with. Int4 everywhere makes the turn slower, 1054 ms with groups of 128, because it costs more on prompts (75 ms) than it saves when decoding (38 ms). The fastest solution within $\Delta\mathcal{L} \leq 0.1$ takes 1004 ms (1.3% faster). The fastest on the front, 995 ms, keeps about 60% of the MLP neurons and costs 0.74 nats, which raises WikiText-2 perplexity from 14.3 to 49.5.

10.3 Stored size against GPU memory

The front spans 158–947 MB of stored size (Figure 12), against 988 MB uncompressed. Table 16 compares the smallest stored solution within each loss budget with the deployed weights of that same solution and with the smallest deployed weights of the three-objective search of Section 8. For int4 models the stored size is only 3–7% below the deployed size, so the last column is a fair rough reference. Sharing nearly halves the stored size at $\Delta\mathcal{L} \leq 0.02$, saves about 27% at 0.05 nats and about 20% at 0.5–1 nats, and adds little around 0.1–0.2 nats, where int4 everywhere is already

	bf16		int4, group 128	
Part of the turn	ms	share	ms	share
Attention over the KV cache, prompt tokens	567	56%	567	54%
Attention over the KV cache, generated tokens	102	10%	102	10%
Fixed parts, norms, activations, residuals	212	21%	212	20%
Attention and MLP products, prompt tokens	16	2%	91	9%
Attention and MLP products, generated tokens	107	10%	74	7%
Output head, generated tokens	13	1%	8	1%
Total	1018		1054	

Table 15: Predicted GPU time of one agent's turn (Eq. 11) of Qwen2.5-0.5B-Instruct on the A100, with 32,000 tokens of context.

$\Delta\mathcal{L}$ up to (nats)	Mix: smallest stored (MB)	Same solution, deployed (MB)	3 obj.: smallest deployed (MB)
0.02	374	966	725
0.05	312	873	430
0.1	272	423	299
0.2	235	417	263
0.5	187	793	228
1.0	158	654	202

Table 16: Stored size of the mixed search against deployed weights, over all evaluated solutions, without the KV cache. The uncompressed model takes 988 MB in both columns.

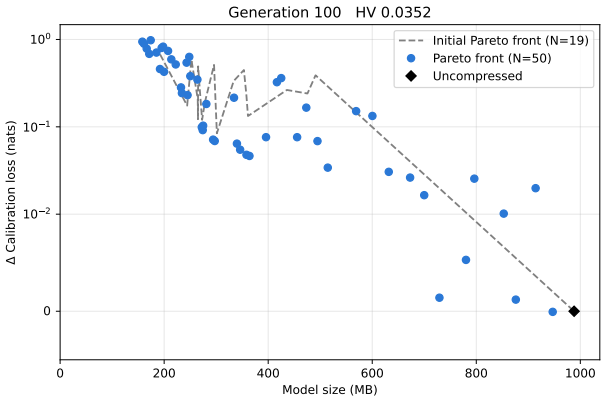
244–301 MB stored. It does so at a cost in GPU memory: the smallest stored solutions keep large matrices in bf16 when running.
GPU memory is limited by the KV cache in this workload: it takes 1573 of the 2561 MB deployed for 4 agents, and every solution fits 180–182 agents on 80 GB. On this model, weight compression therefore does not change how many agents share a GPU.

10.4 What the search chose

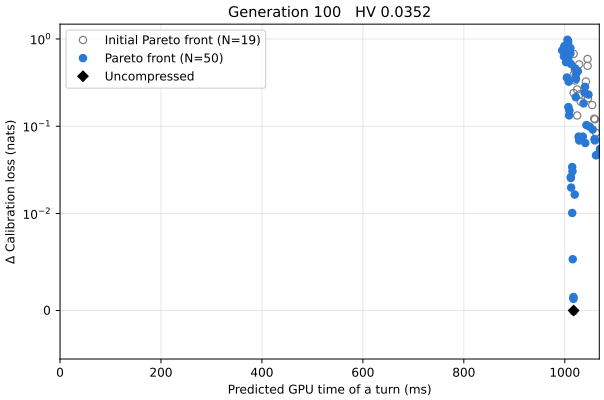
Over the non-dominated solutions among all evaluations (426 within $\Delta\mathcal{L} \leq 1$), 99% use int4 and sharing together, so the mixture does get used. Sharing goes mostly to the MLP (74% of the up genes, 62% of the down genes), which are the largest matrices and the ones on which int4 costs most on prompts; attention projections mostly take int4 with groups of 128 (52%). The extra sharing genes are almost unused: where sharing is used, 97% of the solutions have no clipping ($c = 1$), 98% no pruning ($\beta = 0$) and 98% no merge ($\rho = 1$). The grid level k does vary, from $k = 17$ (37%) to $k = 257$ (11%). Pruning and merging therefore add little on this model once formats and MLP width are available.

10.5 Test perplexity

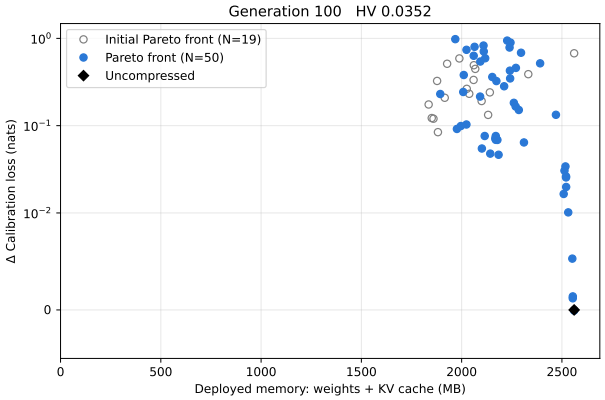
Table 17 and Figure 13 give six points of the final front on 32 windows each. Low-loss points hold up on held-out text (0.055 nats on calibration, 0.081 on WikiText-2 and 0.066 on C4), and high-loss points are optimistic (0.740 becomes 1.24 and 1.08), so the calibration batch of 8 windows overstates what MLP pruning keeps.



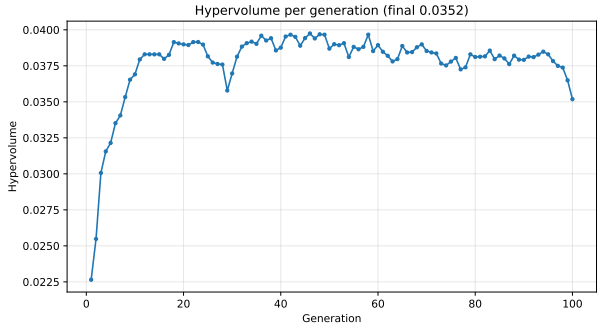
(a) Stored size (objective).



(b) GPU time of a turn (objective).



(c) Deployed memory, 4 agents (reported).

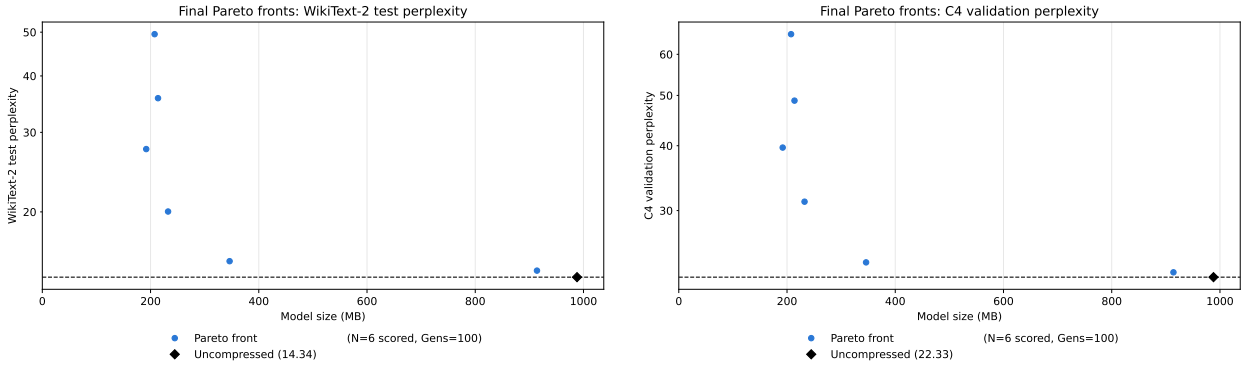


(d) Hypervolume per generation.

Figure 12: Mixed search: final front (50 solutions), initial front and the uncompressed model. A front in three objectives is not ordered along any one axis, so it does not form a curve in these projections.

Turn (ms)	Stored (MB)	Deployed, 4 agents (MB)	$\Delta\mathcal{L}$ (nats)	WikiText-2	C4
1018	988	2561	0	14.34	22.33
1013	914	2521	0.020	14.82	22.81
1070	346	2101	0.055	15.56	23.84
1040	232	2212	0.283	20.04	31.21
1021	192	2270	0.459	27.55	39.67
1007	214	2117	0.592	35.71	48.86
995	207	2024	0.740	49.49	65.61

Table 17: Test perplexity (32 windows of 2048 tokens each) of the uncompressed model and of points of the final front of the mixed search.



(a) WikiText-2 test.

(b) C4 validation.

Figure 13: Test perplexity of the scored points of the mixed search against stored size.

10.6 Search behavior

The hypervolume peaks at generation 46 (0.0397) and ends lower (0.0352), since NSGA-II keeps only 50 solutions of a front that has hundreds, spread over three objectives. Over all evaluations it is 0.0438. This loses good solutions: within $\Delta\mathcal{L} \leq 0.02$ the smallest stored size is 374 MB over all evaluations but 700 MB in the final population. Test scoring sees only the final population, so the best low-loss solutions of Table 16 have a calibration loss and no test score.

10.7 Limitations

- **Prediction only.** This search has no validation step: the turn times are not checked against a real run, and the deployed models with shared matrices are checked only on a small random model. Eager execution adds overhead that the turn time leaves out, as in Tables 10 and 12.
- **The model is not the target.** Qwen2.5-0.5B weights are small next to its KV cache at this context. On the shapes of Qwen2.5-32B the same A100 fits no agent in bf16 (65.5 GB of weights) and 6 with int4 groups of 128 (17.4 GB), measured by a separate kernel probe, so the memory conclusion above does not carry over. Its turn-time ranking does not change: at batch 4 the matrix products with int4 save 3.9 ms per generated token and cost 3.3 ms more per prompt token, about 5.5 s more per turn of this workload.
- **The KV cache is not a decision.** The largest part of both the time and the memory is outside the search space.
- **Other limits.** One seed, one workload, sublayer removal off, and perplexity only: none of the A100 solutions has been run on the downstream benchmarks or on agent tasks.

11 Discussion

When the third objective should matter. Latency and memory are nearly collinear on Qwen2.5-0.5B for two reasons: its embedding is tied to the output head, so compressing it shortens the decode step as well, and its grouped-query attention keeps the KV cache at 12.6 MB. nanochat d34 differs on both counts. Its untied embedding of 65536×2176 weights (285 MB in bf16) costs memory but no decode time, since a lookup reads one row; a two-objective search has no reason to compress it, while a three-objective search does. Its multi-head attention keeps 296 KB of keys

and values per token (303 MB at 1024 tokens), which grows with the context and the batch. A comparison on d34 would therefore show whether the third objective earns its cost.

GPU time against wall-clock time. The latency objective is GPU time, which is what a runtime with CUDA graphs or a serving engine pays. Eager PyTorch at batch 1 is limited by the CPU for a model of this size: the deployed models of Table 5 take 23–35 ms per forward in eager mode against 2.8–6.8 ms of GPU time, and the same uncompressed model measured 28.0 and 34.6 ms in two processes.

Limitations.

- One seed per scenario: the comparisons of Sections 7 and 9 are indicative only.
- The kernel menu is that of this machine: no LUT, 2:4 or fused int8 kernels, and no FlashAttention.
- The hardware-aware searches model decoding at batch 1 only, and the tinygemm int4 kernel is 3–5× slower than bf16 on long prompts on the RTX 3060 (Table 2); the search of Section 10 models prompts as well.
- Structurally pruned models are not fine-tuned afterwards, which makes structural pruning look worse than it is after a short recovery training [13].
- The loss is measured on 8 calibration windows, and the local test sets on 32 windows each.
- The latency table belongs to one GPU and one software stack and must be measured on an otherwise idle GPU.

12 Future work

- **Lookup-table kernels.** With a LUT kernel such as FLUTE [9] or the kernels of SqueezeLLM [11], non-uniform levels cost nothing at runtime. The per-row scale and global codebook of UBQ-ECSQ map directly onto such a kernel, with codebooks of exactly 2^b entries at 2, 3 or 4 bits.
- **Sparsity and fused kernels on Linux.** 2:4 sparsity [16] (and int4 with 2:4 in Sparse-Marlin [8]), fused int8 activations for prompt-heavy workloads [1, 23], and KV-cache quantization for long contexts and batches.
- **Per-layer genes.** Decisions per layer instead of per band, seeded by a separable sensitivity model solved as a knapsack problem under a latency budget, as in HAWQ-V3 [24], with the same model screening offspring.
- **KV-cache size and precision as decisions,** since the cache dominates the turn time and the memory of Section 10.
- **Statistics.** Several seeds per scenario, compared by the mean final hypervolume and a Mann–Whitney U test.
- **Recovery fine-tuning** of the structurally pruned points of the final front.

Matrix ($N \times K$)	Role	bf16	int4 g32	int4 g64	int4 g128
896×896 + bias	q	9.6	5.0	4.7	4.6
128×896 + bias	k, v	3.6	2.5	2.5	2.5
896×896	o	10.8	3.9	3.6	3.5
4864×896	gate, up	40.7	13.8	12.6	11.9
896×4864	down	46.1	13.9	12.7	12.0
3584×896	gate, up ($\kappa = \frac{3}{4}$)	30.7	10.6	9.8	9.3
896×3584	down ($\kappa = \frac{3}{4}$)	34.6	10.6	9.7	9.3
2432×896	gate, up ($\kappa = \frac{1}{2}$)	21.6	7.6	7.0	6.7
896×2432	down ($\kappa = \frac{1}{2}$)	24.6	7.6	7.0	6.7
151936×896	output head	1201.5	380.0	340.9	321.3

Table 18: GPU time (μs) of one product with a single token. Attention over 1024 cached tokens takes 19.9 μs per layer (grouped implementation; SDPA: 130.6 μs).

A Measured matrix-product times for Qwen2.5-0.5B

B Sublayer removal ranking for Qwen2.5-0.5B

Loss increase (nats) on 4 calibration windows when one sublayer alone is removed; the search removes sublayers in this order.

- **Attention**, first eight: layer 18 (0.027), 13 (0.032), 6 (0.035), 19 (0.038), 7 (0.040), 22 (0.042), 9 (0.043), 11 (0.060); last: 23 (0.342), 0 (2.37), 1 (5.12).
- **MLP**, first eight: layer 12 (0.056), 10 (0.057), 11 (0.069), 9 (0.079), 13 (0.082), 14 (0.084), 8 (0.095), 18 (0.119); last: 23 (0.537), 2 (2.08), 0 (7.94).

The middle layers are the most redundant and the first and last ones the least, in line with ShortGPT [14].

References

[1] Saleh Ashkboos, Amirkeivan Mohtashami, Maximilian L. Croci, Bo Li, Pashmina Cameron, Martin Jaggi, Dan Alistarh, Torsten Hoefler, and James Hensman. QuaRot: Outlier-free 4-bit inference in rotated LLMs. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

[2] Julian Blank and Kalyanmoy Deb. pymoo: Multi-objective optimization in Python. *IEEE Access*, 8:89497–89509, 2020.

[3] Dimo Brockhoff and Eckart Zitzler. Objective reduction in evolutionary multiobjective optimization: theory and applications. *Evolutionary Computation*, 17(2):135–166, 2009.

[4] Han Cai, Ligeng Zhu, and Song Han. ProxylessNAS: Direct neural architecture search on target task and hardware. In *International Conference on Learning Representations (ICLR)*, 2019.

[5] Tri Dao. FlashAttention-2: Faster attention with better parallelism and work partitioning. In *International Conference on Learning Representations (ICLR)*, 2024.

[6] Kalyanmoy Deb, Amrit Pratap, Sameer Agarwal, and T. Meyarivan. A fast and elitist multiobjective genetic algorithm: NSGA-II. *IEEE Transactions on Evolutionary Computation*, 6(2):182–197, 2002.

[7] Elias Frantar, Saleh Ashkboos, Torsten Hoefler, and Dan Alistarh. GPTQ: Accurate post-training quantization for generative pre-trained transformers. In *International Conference on Learning Representations (ICLR)*, 2023.

[8] Elias Frantar, Roberto L. Castro, Jiale Chen, Torsten Hoefler, and Dan Alistarh. MARLIN: Mixed-precision auto-regressive parallel inference on large language models. In *ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP)*, 2025.

[9] Han Guo, William Brandon, Radostin Cholakov, Jonathan Ragan-Kelley, Eric P. Xing, and Yoon Kim. Fast matrix multiplications for lookup table-quantized LLMs. In *Findings of the Association for Computational Linguistics: EMNLP*, 2024.

[10] Andrej Karpathy. nanochat. `https://github.com/karpathy/nanochat`, 2025.

[11] Sehoon Kim, Coleman Hooper, Amir Gholami, Zhen Dong, Xiuyu Li, Sheng Shen, Michael W. Mahoney, and Kurt Keutzer. SqueezeLLM: Dense-and-sparse quantization. In *International Conference on Machine Learning (ICML)*, 2024.

[12] Ji Lin, Jiaming Tang, Haotian Tang, Shang Yang, Wei-Ming Chen, Wei-Chen Wang, Guangxuan Xiao, Xingyu Dang, Chuang Gan, and Song Han. AWQ: Activation-aware weight quantization for on-device LLM compression and acceleration. In *Proceedings of Machine Learning and Systems (MLSys)*, 2024.

[13] Xinyin Ma, Gongfan Fang, and Xinchao Wang. LLM-Pruner: On the structural pruning of large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

[14] Xin Men, Mingyu Xu, Qingyu Zhang, Bingning Wang, Hongyu Lin, Yaojie Lu, Xianpei Han, and Weipeng Chen. ShortGPT: Layers in large language models are more redundant than you expect. *arXiv preprint arXiv:2403.03853*, 2024.

[15] Stephen Merity, Caiming Xiong, James Bradbury, and Richard Socher. Pointer sentinel mixture models. In *International Conference on Learning Representations (ICLR)*, 2017.

[16] Asit Mishra, Jorge Albericio Latorre, Jeff Pool, Darko Stosic, Dusan Stosic, Ganesh Venkatesh, Chong Yu, and Paulius Micikevicius. Accelerating sparse deep neural networks. *arXiv preprint arXiv:2104.08378*, 2021.

[17] PyTorch Team. torchao: PyTorch native quantization and sparsity for training and inference. `https://github.com/pytorch/ao`, 2024.

[18] Qwen Team. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.

[19] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21(140):1–67, 2020.

[20] Mingjie Sun, Zhuang Liu, Anna Bair, and J. Zico Kolter. A simple and effective pruning approach for large language models. In *International Conference on Learning Representations (ICLR)*, 2024.

[21] Kuan Wang, Zhijian Liu, Yujun Lin, Ji Lin, and Song Han. HAQ: Hardware-aware automated quantization with mixed precision. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019.

[22] Bichen Wu, Xiaoliang Dai, Peizhao Zhang, Yanghan Wang, Fei Sun, Yiming Wu, Yuandong Tian, Peter Vajda, Yangqing Jia, and Kurt Keutzer. FBNet: Hardware-aware efficient ConvNet design via differentiable neural architecture search. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019.

[23] Guangxuan Xiao, Ji Lin, Mickael Seznec, Hao Wu, Julien Demouth, and Song Han. SmoothQuant: Accurate and efficient post-training quantization for large language models. In *International Conference on Machine Learning (ICML)*, 2023.

[24] Zhewei Yao, Zhen Dong, Zhangcheng Zheng, Amir Gholami, Jiali Yu, Eric Tan, Leyuan Wang, Qijing Huang, Yida Wang, Michael W. Mahoney, and Kurt Keutzer. HAWQ-V3: Dyadic neural network quantization. In *International Conference on Machine Learning (ICML)*, 2021.

[25] Eckart Zitzler and Lothar Thiele. Multiobjective evolutionary algorithms: a comparative case study and the strength Pareto approach. *IEEE Transactions on Evolutionary Computation*, 3 (4):257–271, 1999.